

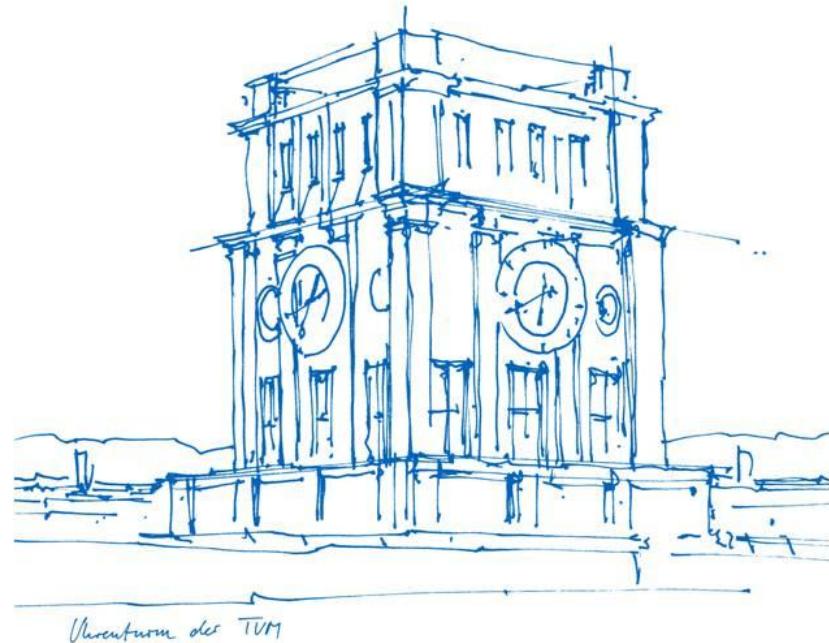
Geometrical Deep Learning on 3D Models: Classification for Additive Manufacturing

Nouhayla Bouziane | Ahmed Ebid | Aditya Sai Srinivas | Felix Bok | Johannes Kiechle

Technical University Munich

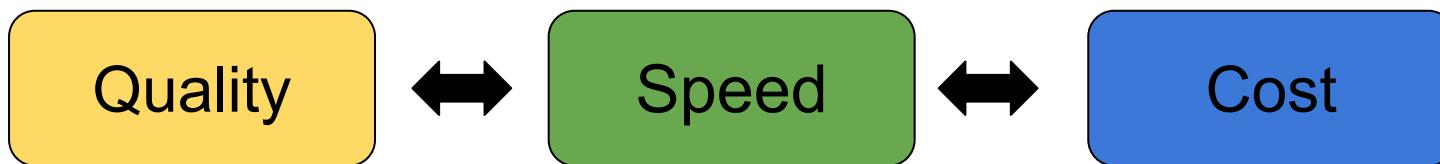
Faculty of Mathematics

22. July 2021



Motivation

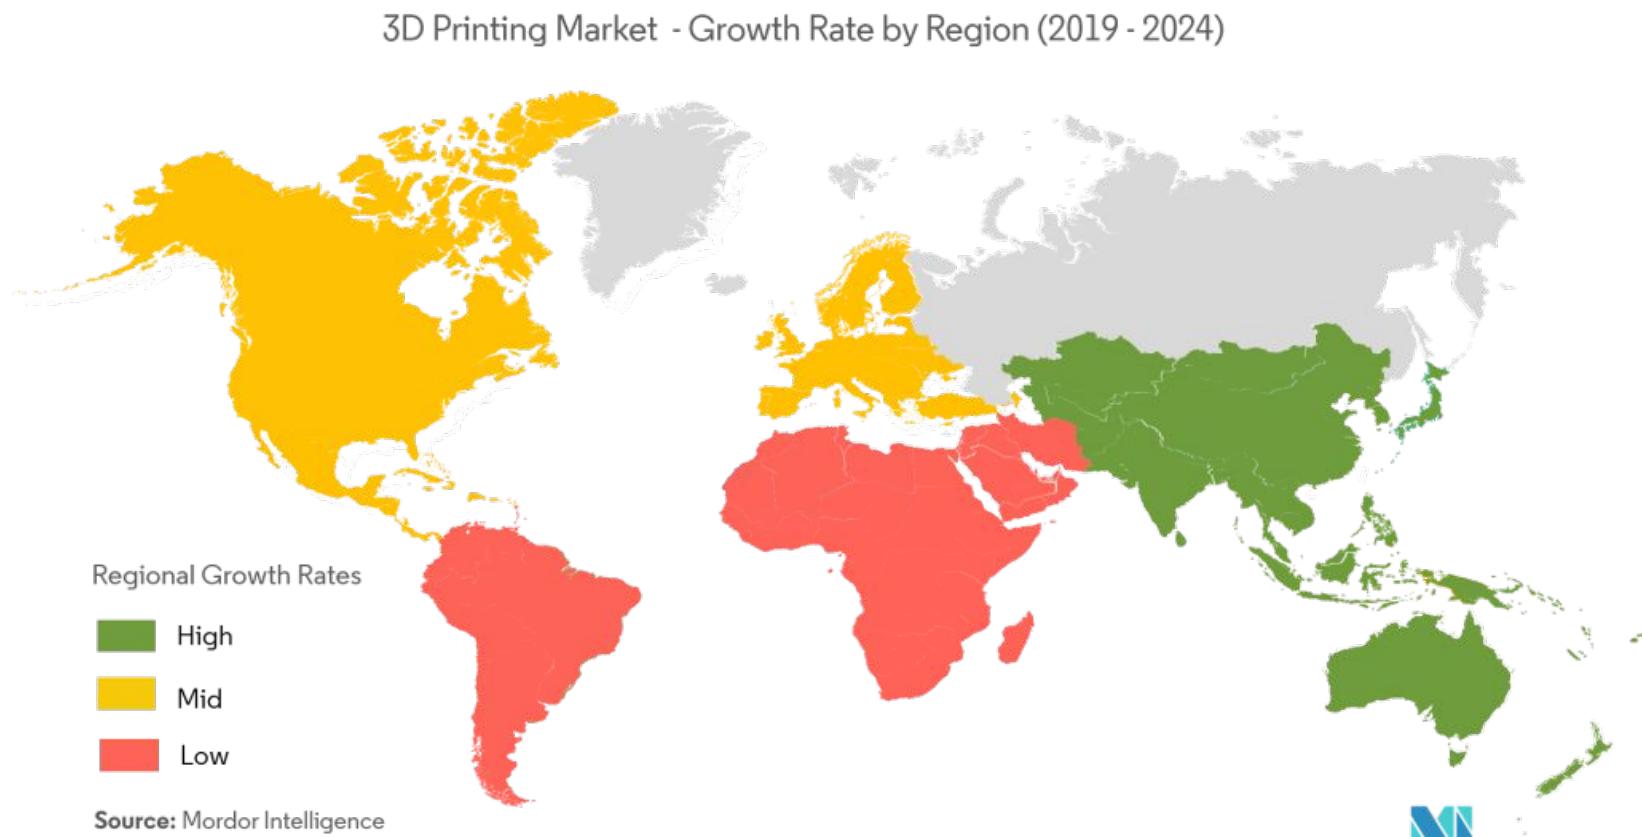
Hard to achieve the right balance:



Additive manufacturing here to transform production processes:

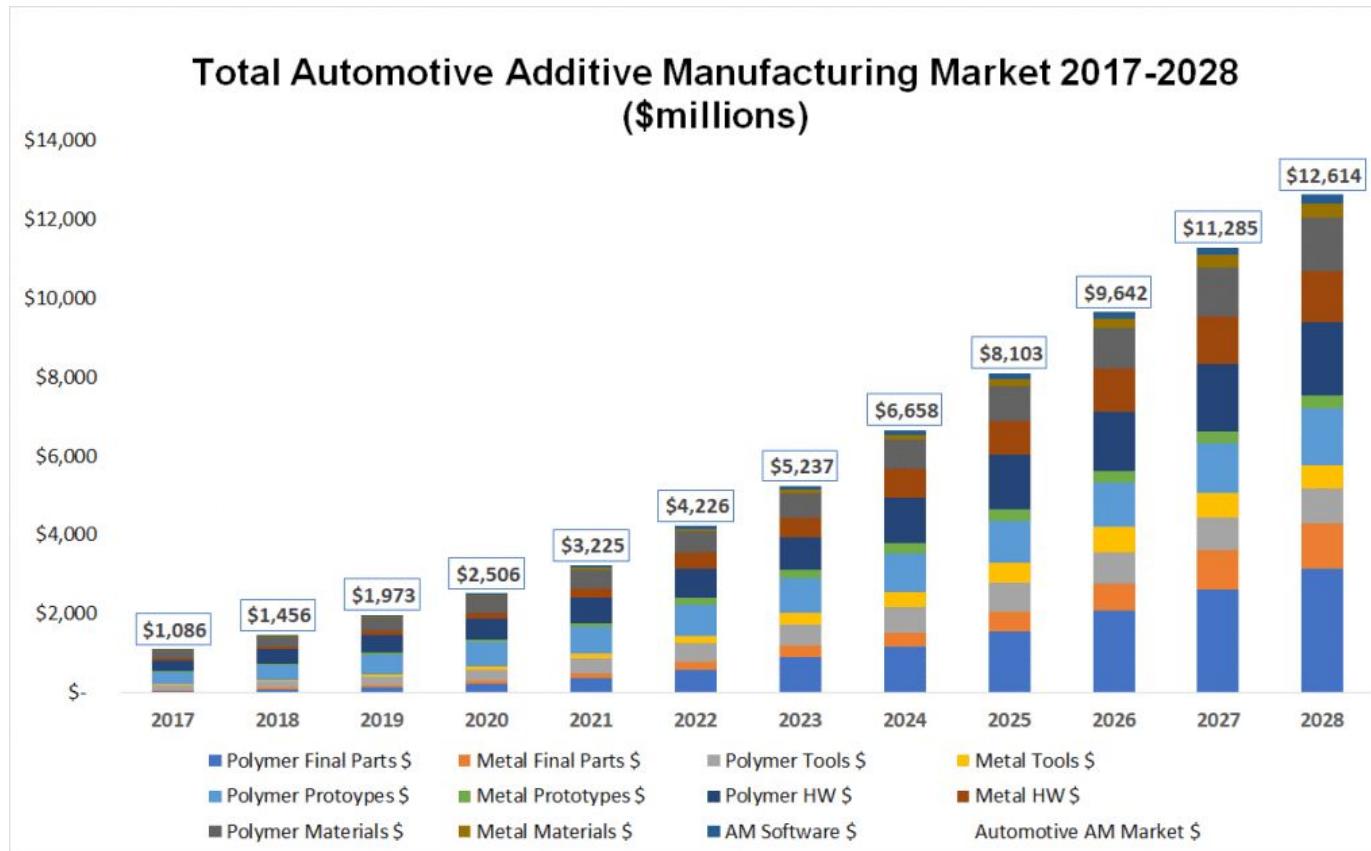
- create design iterations
- enhance quality through cost-effective prototyping
- create specific tooling parts

Motivation

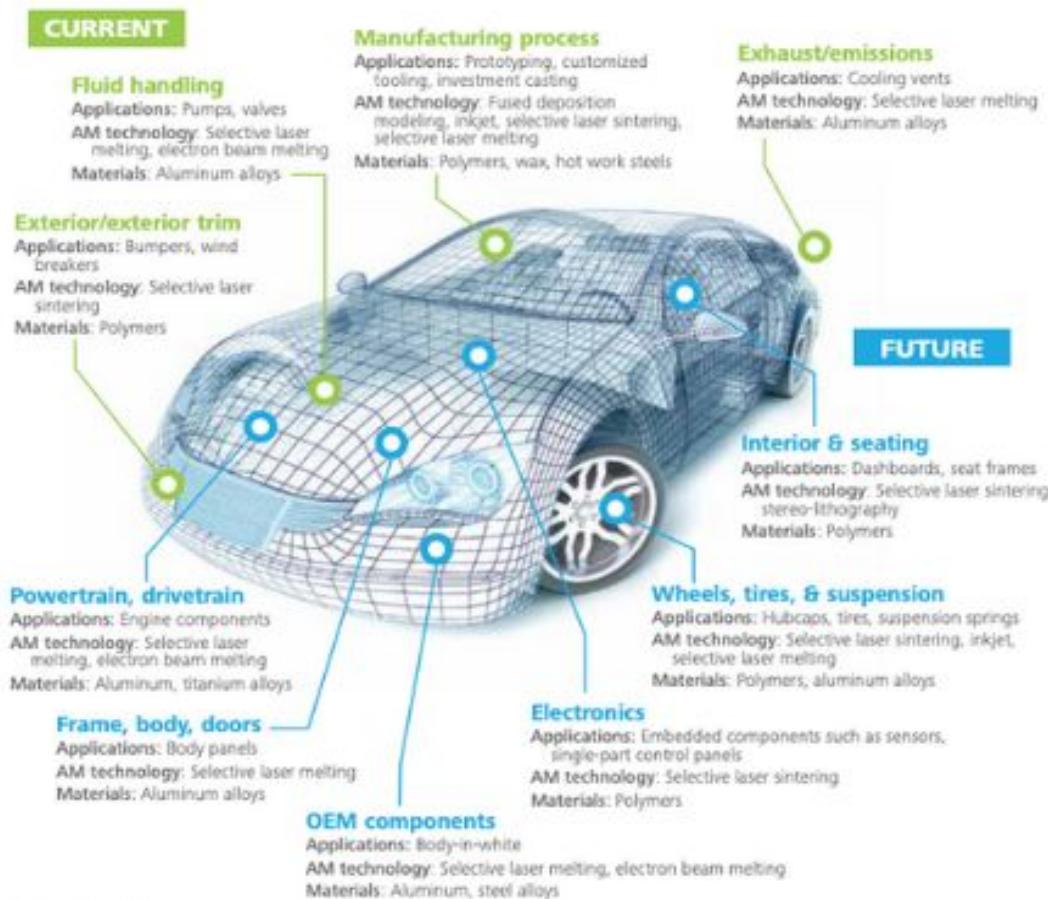


Motivation

By 2028 AM industry revenues would hit **\$12.6** Billion in the automotive Industry



Motivation



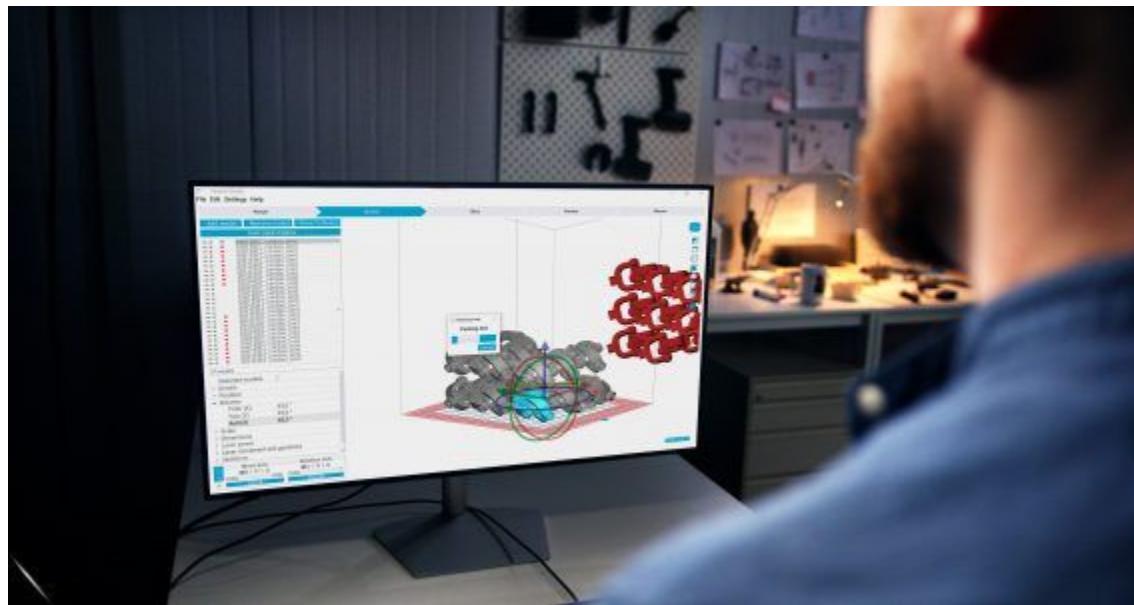
Source: Deloitte analysis.

Applications of AM in the automotive industry

Problem Definition & Project Goal

Requires a lot of human expertise and supervision.

In particular, the process of **identifying whether a 3D model is manufacturable by a given 3D printer** is a **very time-consuming and complex task**.

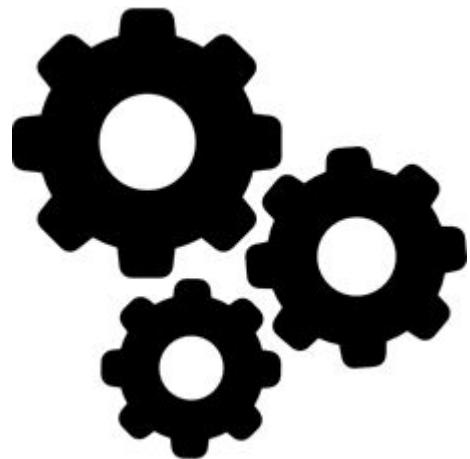


Automate this process using advanced **convolutional neural networks (CNNs)**

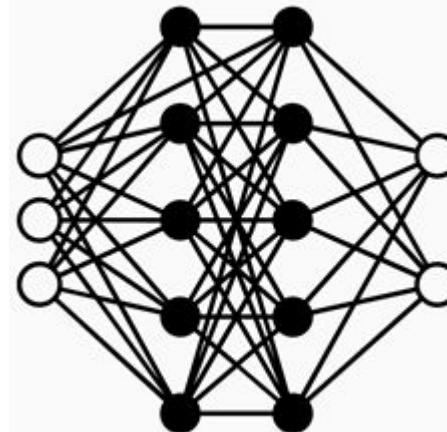
Project Structure

LRZ AI System

Data Generation



Deep Learning

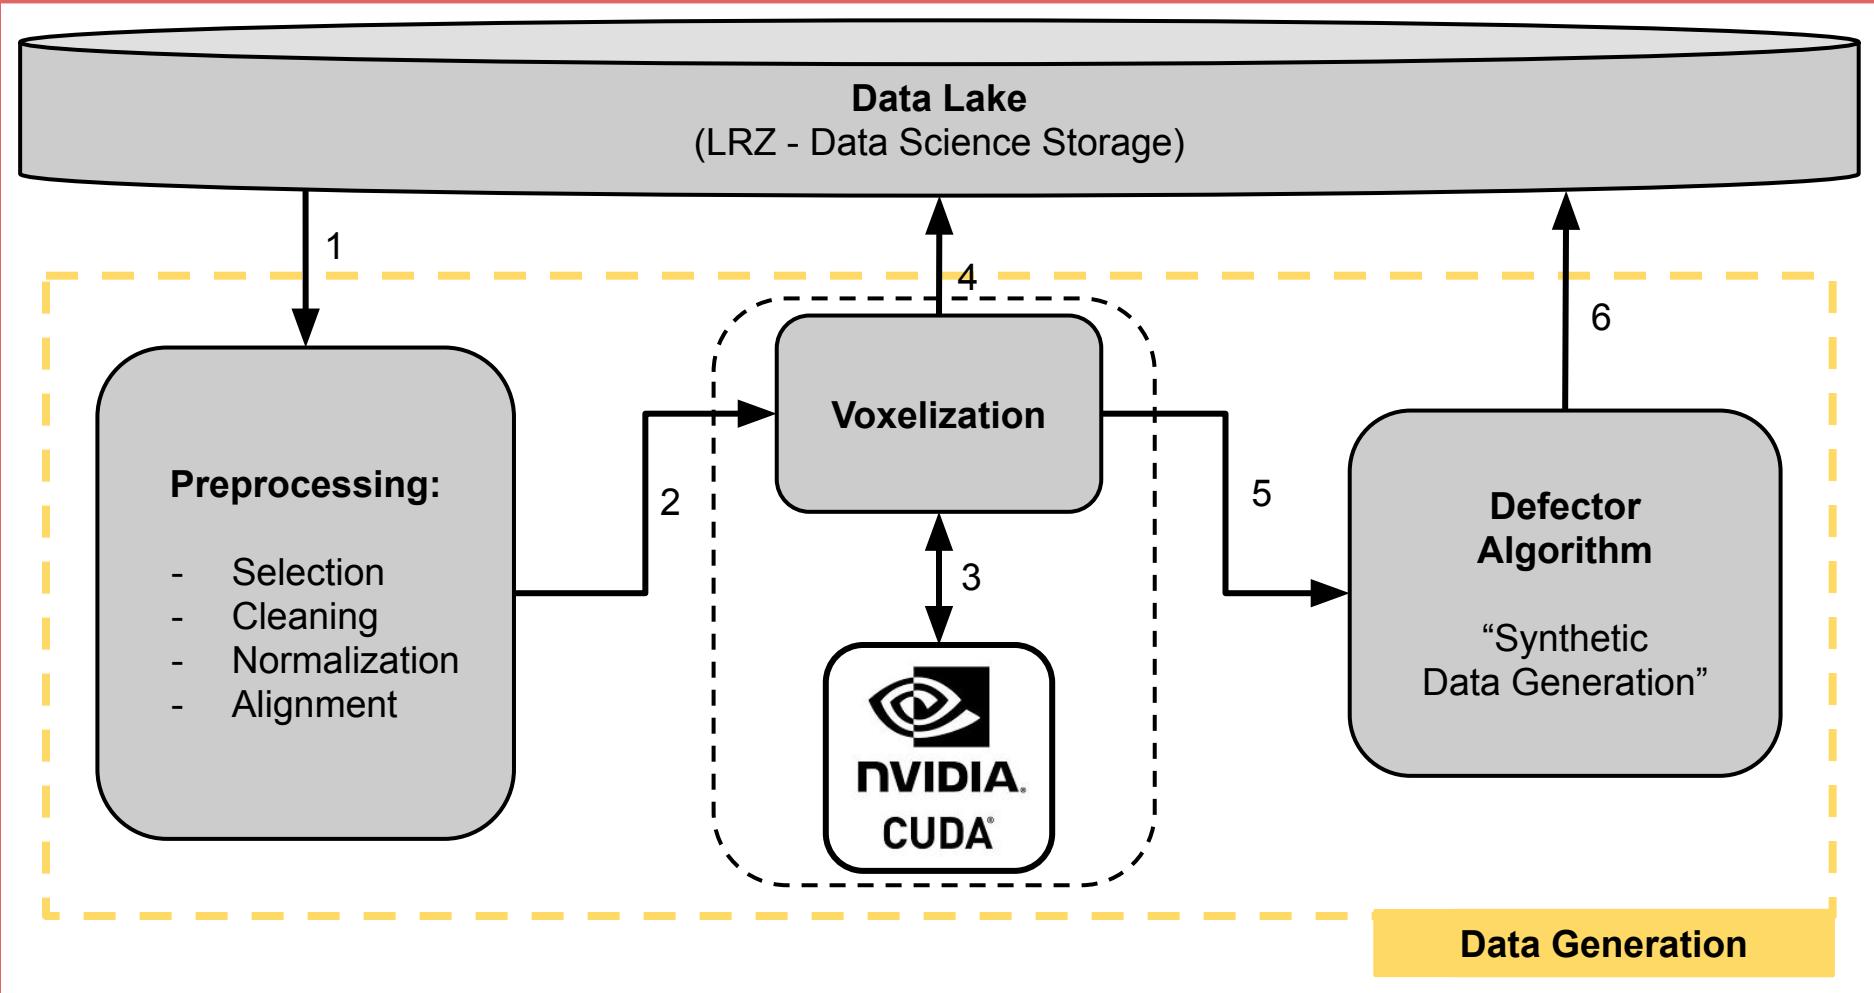


Performance Analysis



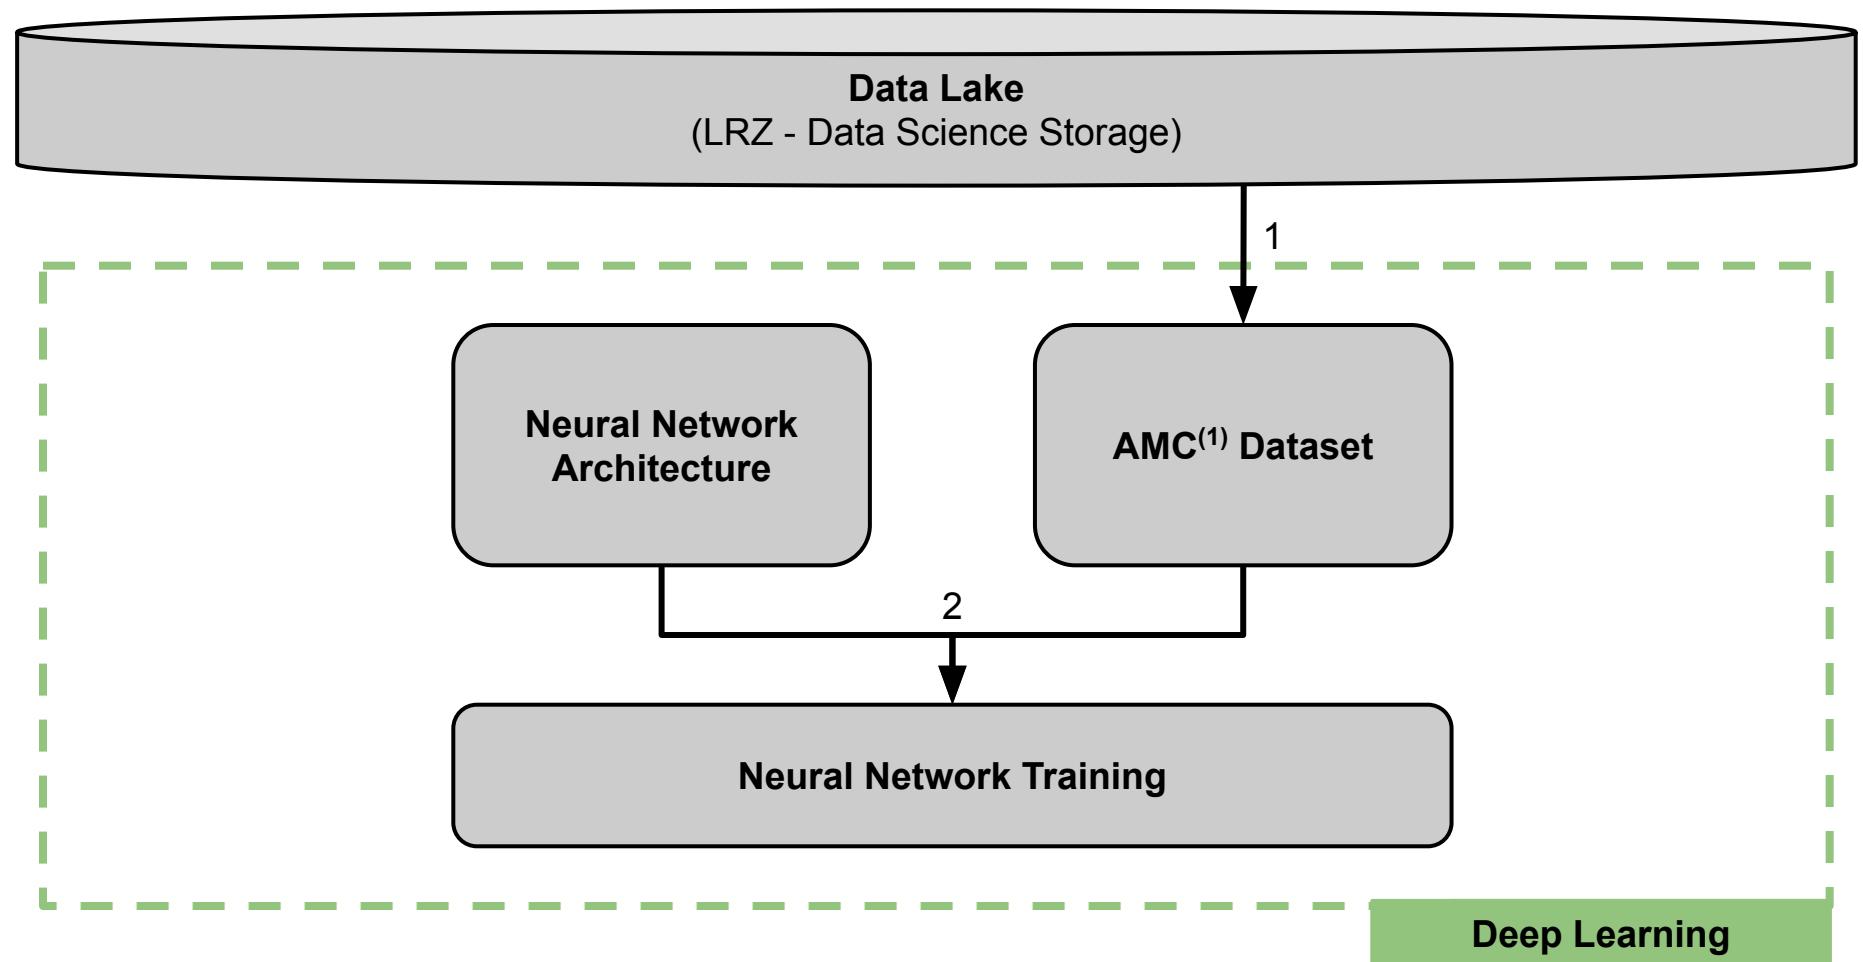
Infrastructure

LRZ AI System



Infrastructure

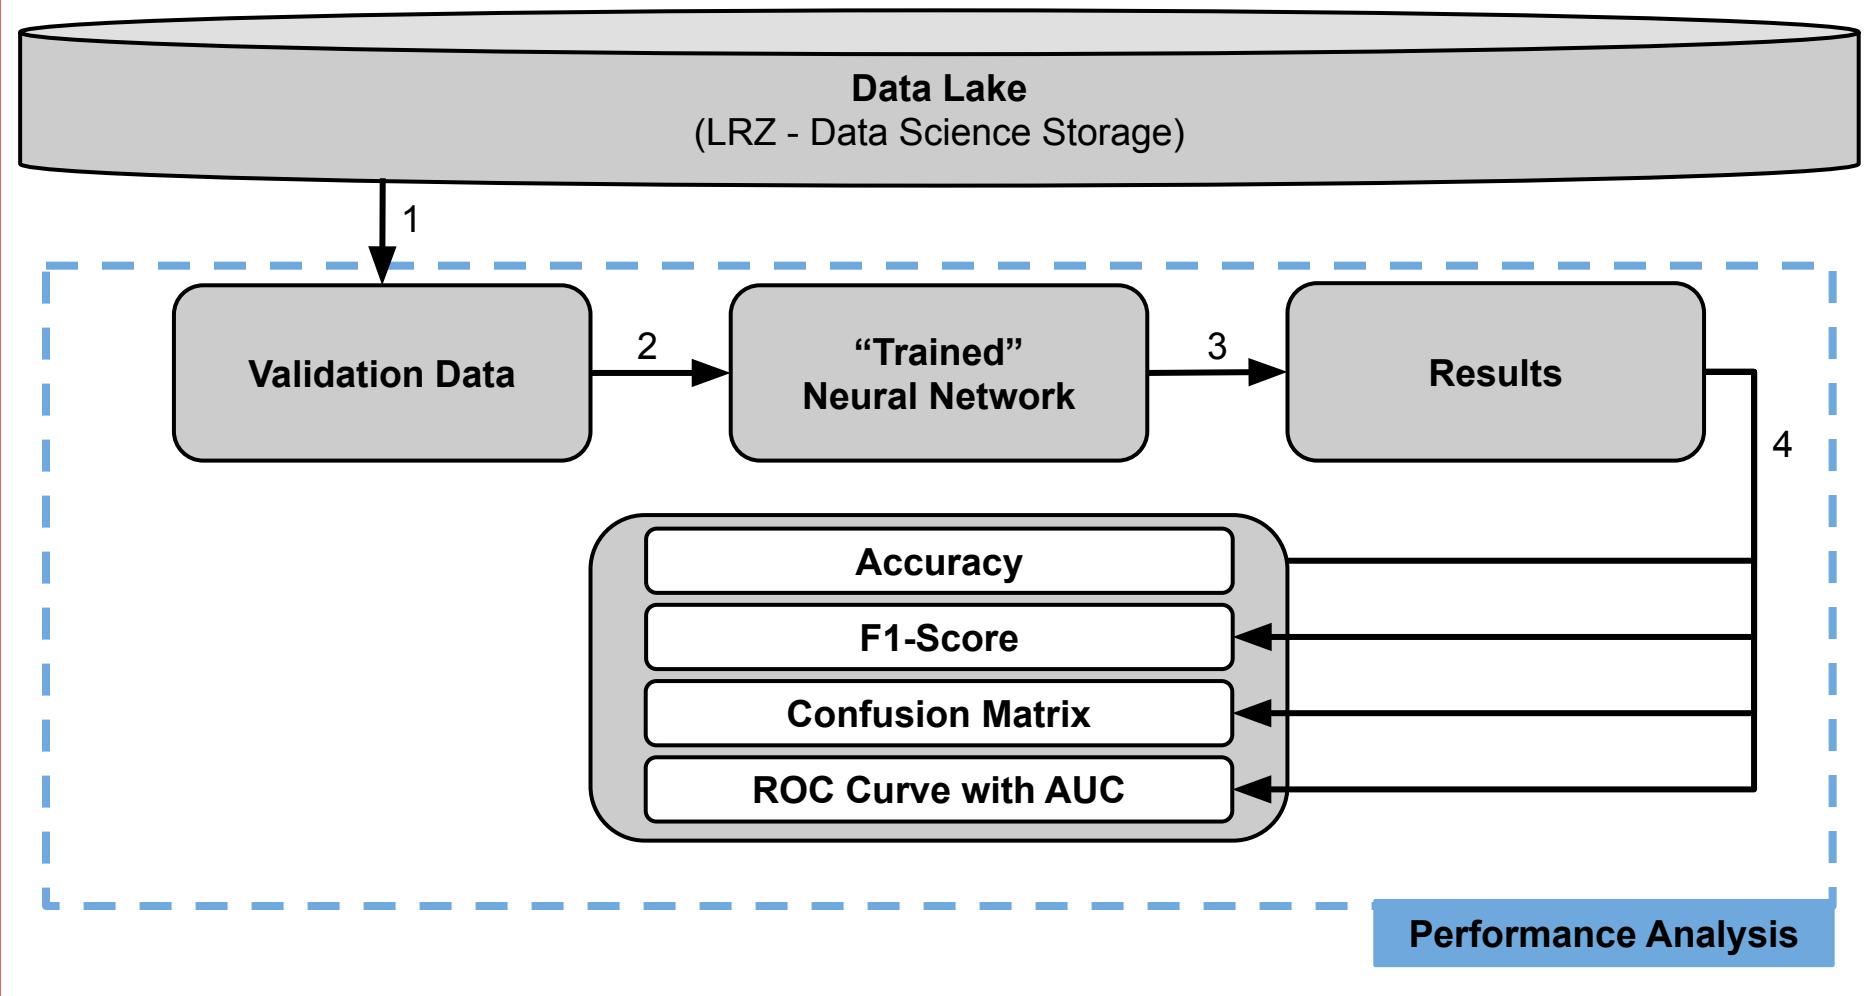
LRZ AI System



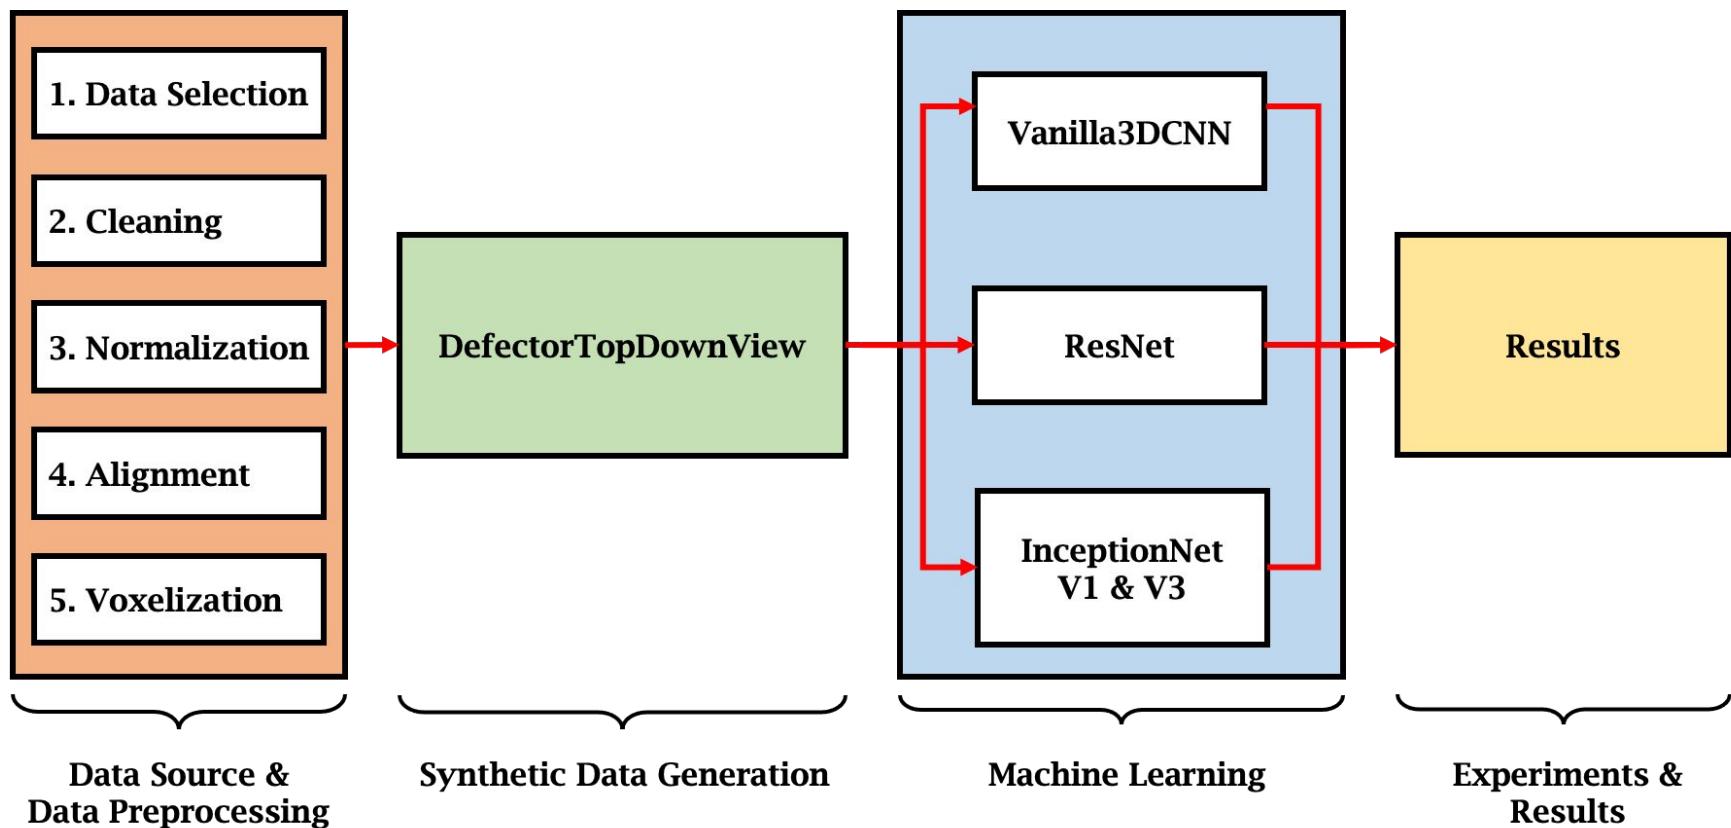
(1) AMC: Additive Manufacturing Classification

Infrastructure

LRZ AI System

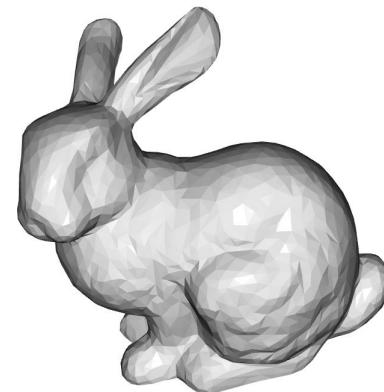


Holistic Project Overview

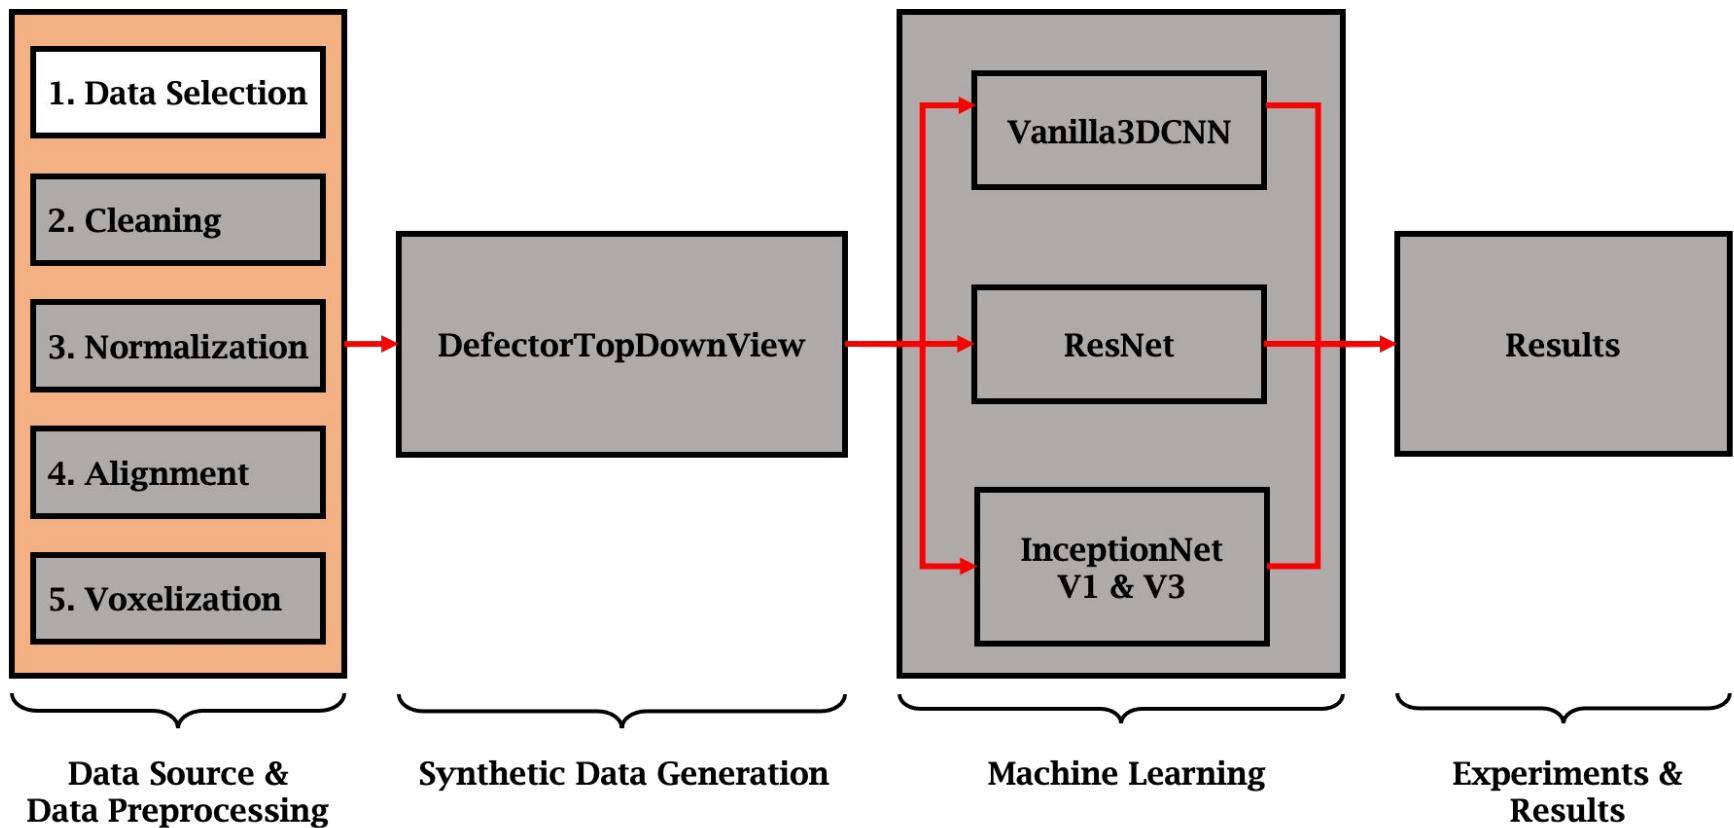


Data Format: Triangular Mesh

- A triangular mesh is a data representation of 3D objects.
- It consists of two main data structures: **Vertices** and **Triangles**.
- **Vertices** is a list of 3D points.
- **Triangles** is a list of triangles, where a triangle connects 3 vertices.

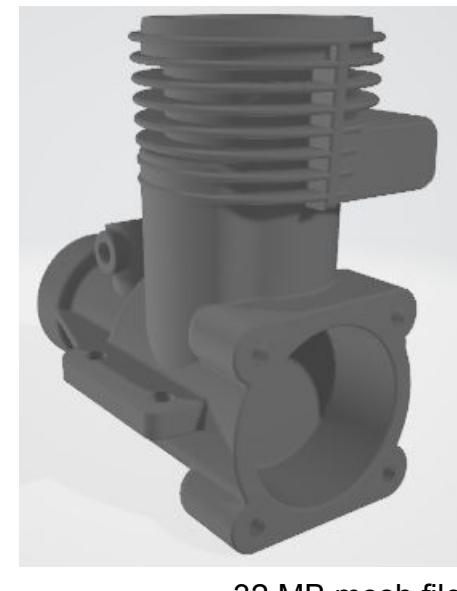
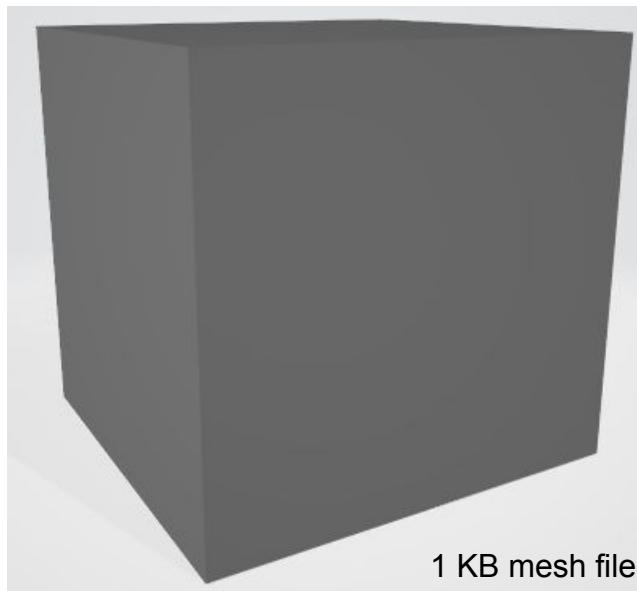


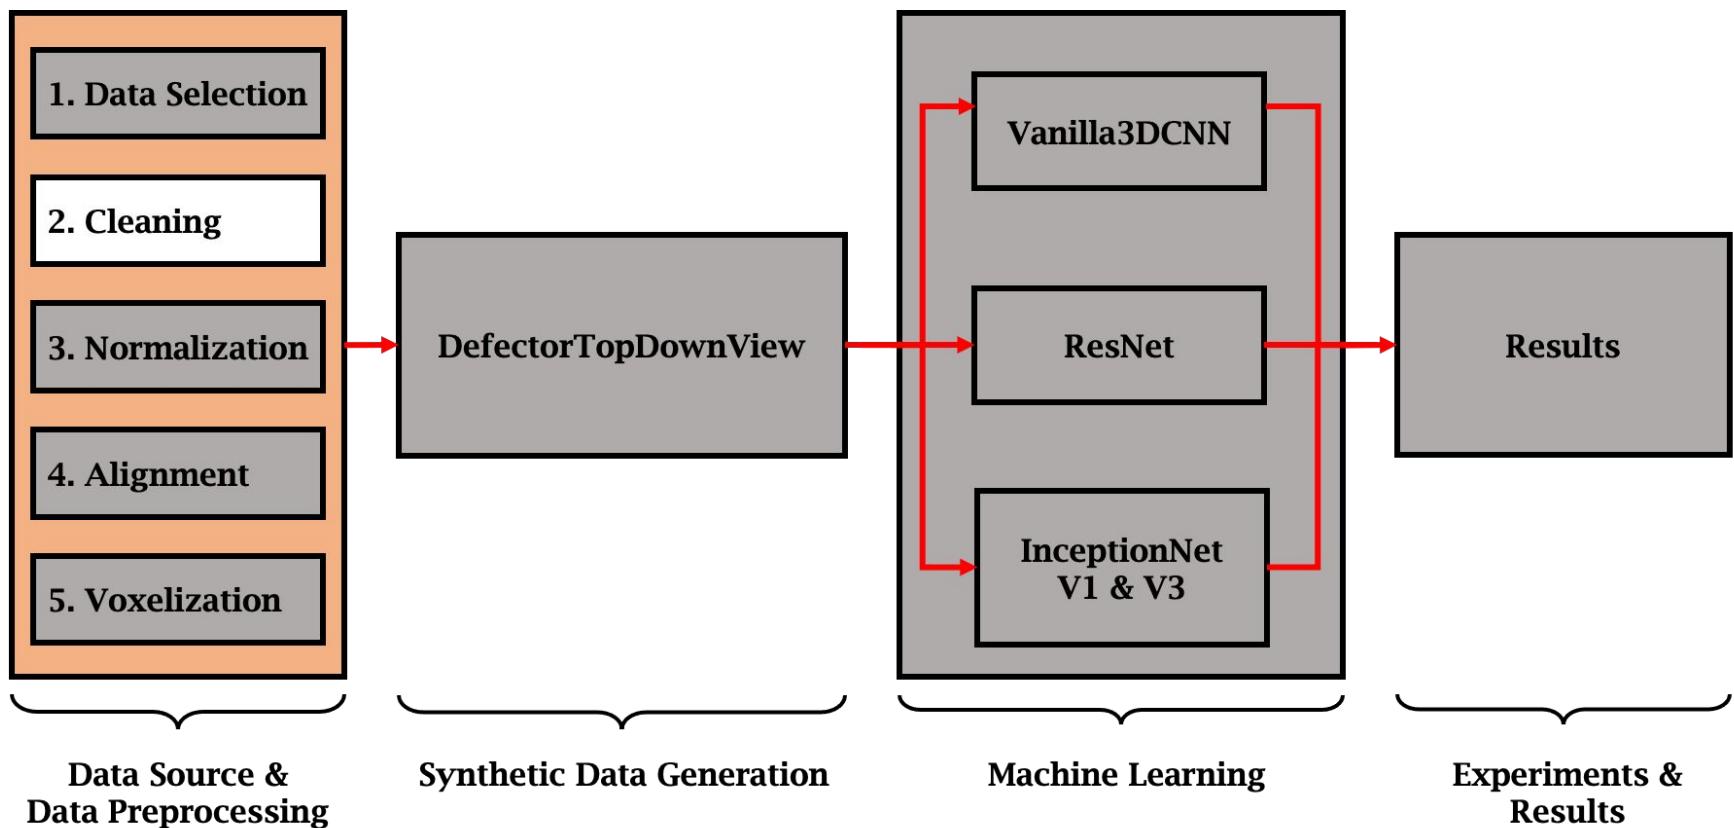
Triangular Mesh Representation



Data Selection

- dataset contains models of varying complexity
- more triangles are needed in meshes of higher complexity
- more triangles translated to bigger file size
- we applied a cutoff value of 25 KB to select models considered for further preprocessing steps



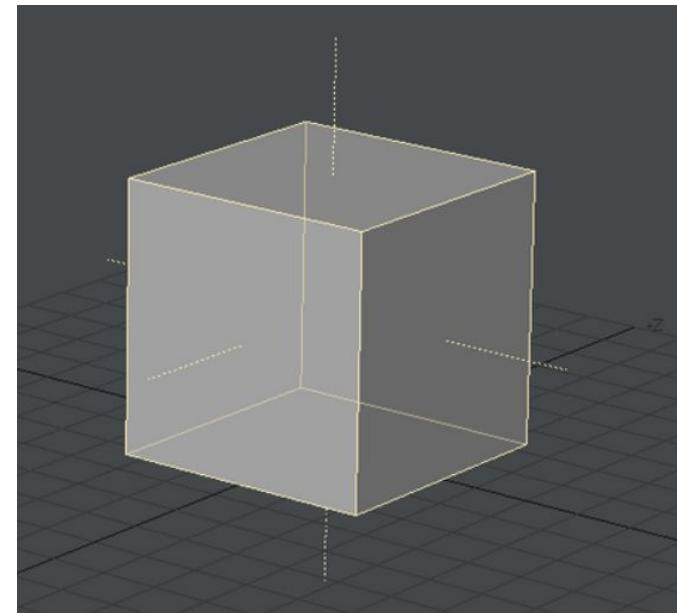


Cleaning

- For a 3D mesh to be 3D-printable, 2 main properties should be satisfied:
 - 1. **Watertightness**: mesh has no holes + normals are facing outwards



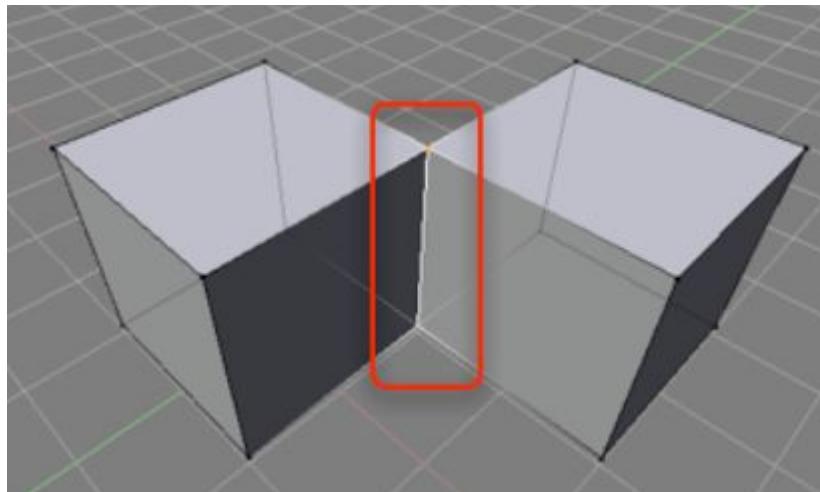
Stanford Bunny Model Bottom View
Invalid: has holes



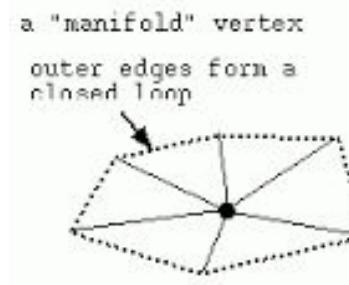
Valid: normals are facing outwards

Cleaning

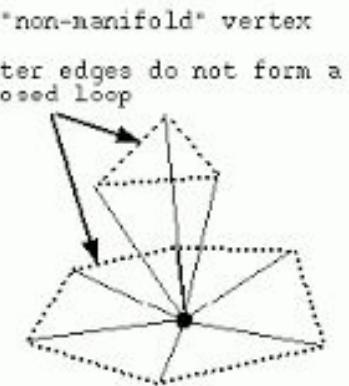
- For a 3D mesh to be 3D-printable, 2 main properties should be satisfied
 - **2. Manifold geometry:** mesh has no edges shared by more than two faces



Invalid: edge shared by 4 faces



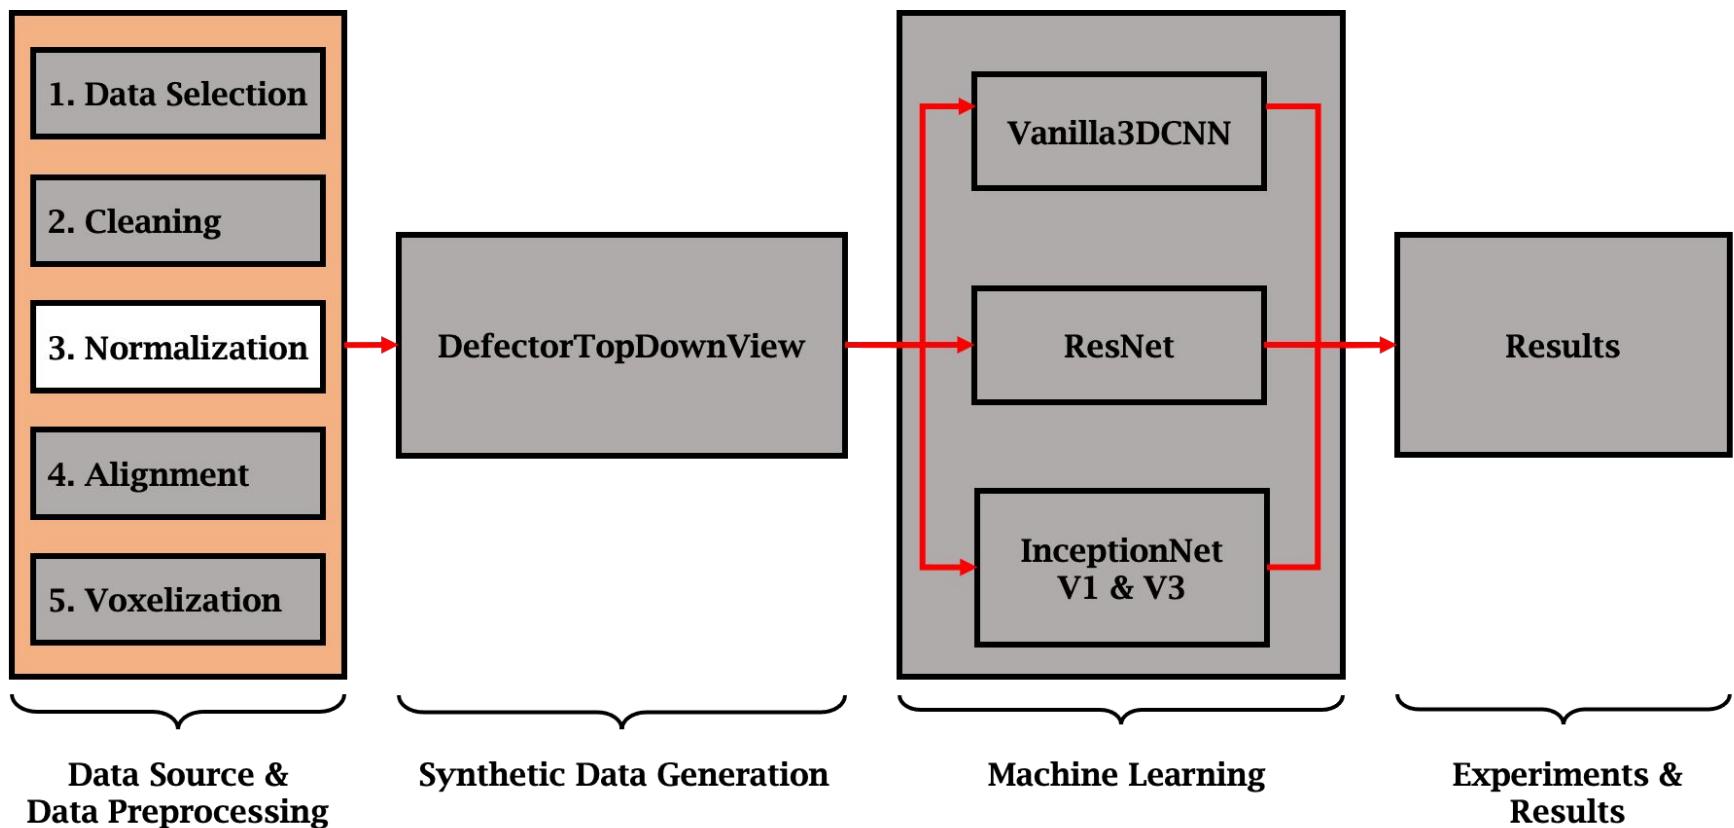
a "manifold" vertex
outer edges form a
closed loop



a "non-manifold" vertex
outer edges do not form a
closed loop

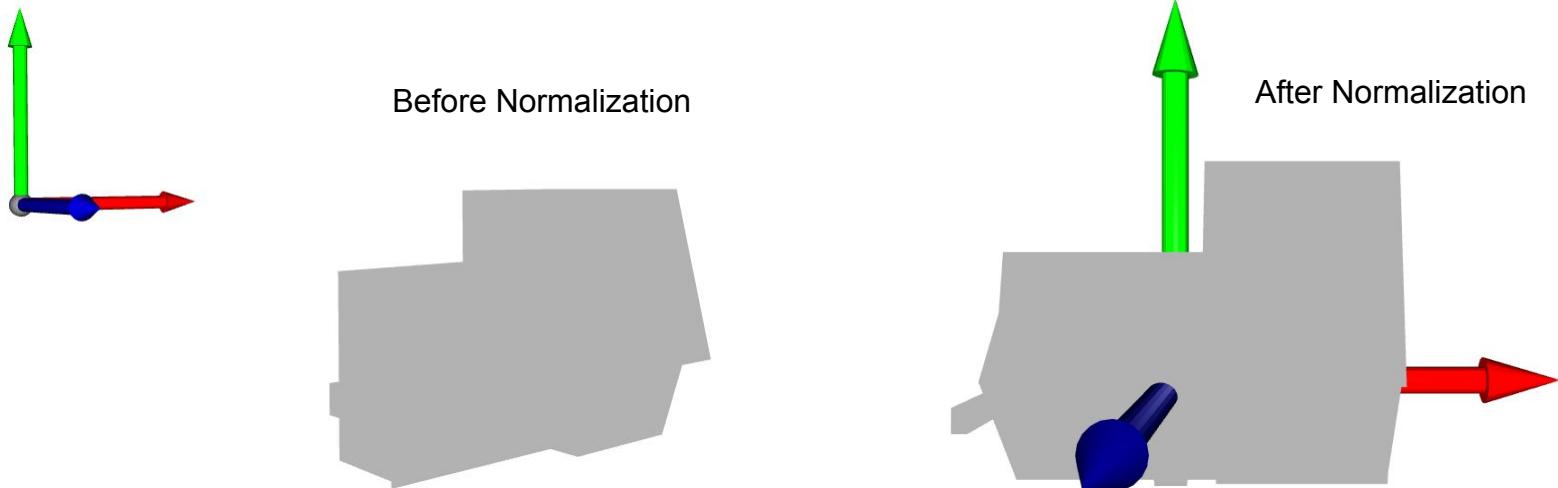
Cleaning

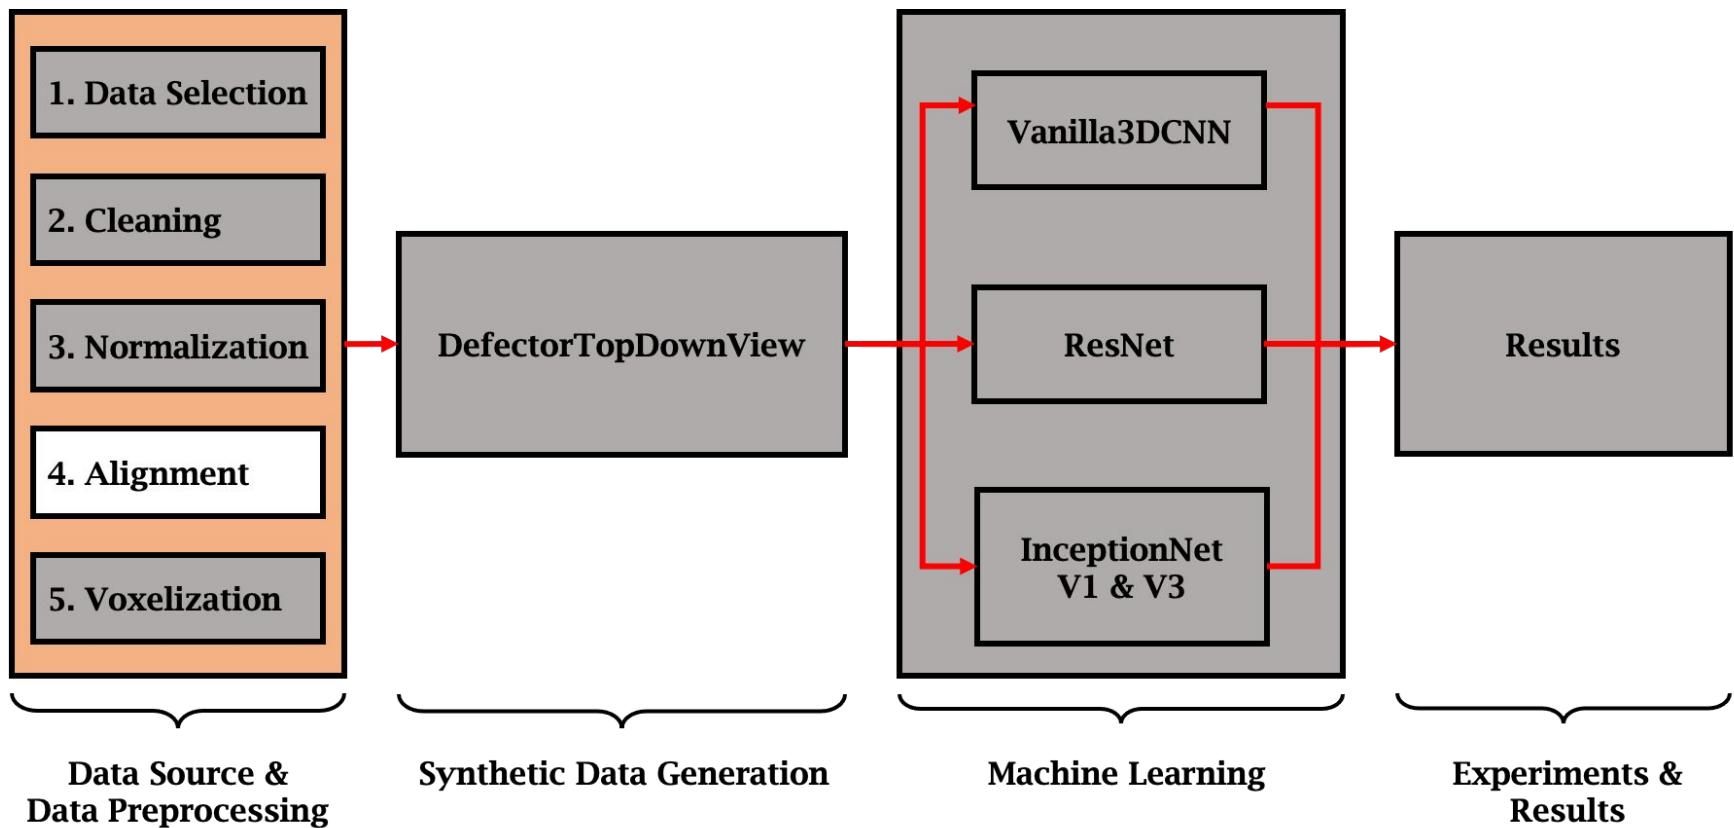
- Mesh cleaning functions were utilized from O3D. [1]
 - **Vertices fixes**
 - Remove vertices that have identical coordinates ($[x1, y1, z1], [x1, y1, z1]$)
 - Remove vertices that are not referenced in any triangle
 - **Edges fixes**
 - Remove non-manifold edges
 - **Triangles fixes**
 - Remove triangles that reference the same three vertices ($[v1, v2, v3], [v2, v1, v3]$)
 - Remove triangles that reference a single vertex multiple times in a single triangle ($[v1, v2, v2]$)



Normalization

- Performed to make sure all vertices of different objects lie in the same range
- **Step 1:** Center the mesh around the origin
 - Find the center of the mesh vertices
 - Translate the mesh vertices to the origin by subtracting the center from all vertices
- **Step 2:** Scale the vertices so that they lie in a [-1, 1] range
 - Divide the mesh vertices by the difference between the maximum bounding point and the minimum bounding point.





Alignment

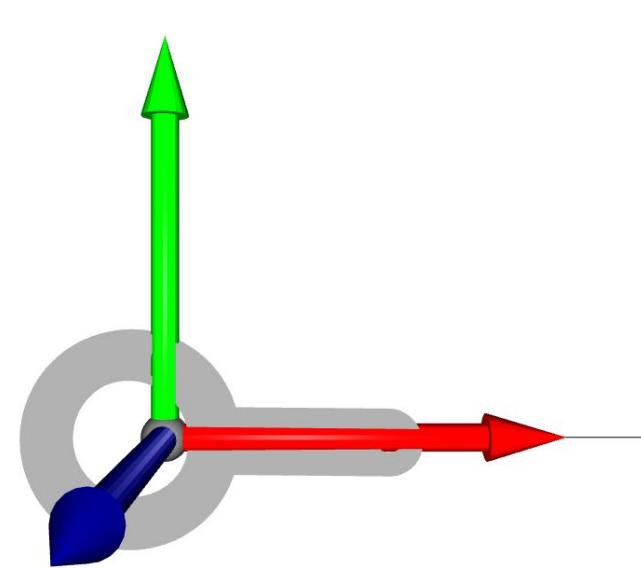
- Performed to make sure all meshes are presented in the same orientation.
- The axis of the minimum Moment Of Inertia (MOI) of a mesh represents the axis around which most of the mass of the object is wrapped. [1]
- Aligning the axis of the minimum MOI with one of the coordinate axes will allow meshes to be presented in the same orientation
- How its done [2] & [3]:
 - **Step 1:** Find the axis of the minimum MOI
 - **Step 2:** Compute a rotation matrix that aligns the axis of the minimum MOI with a coordinate axis
 - **Step 3:** Apply the rotation matrix on the mesh vertices

[1] James Dann and James J. Dann. The People's Physics Book. third edition, 2006.

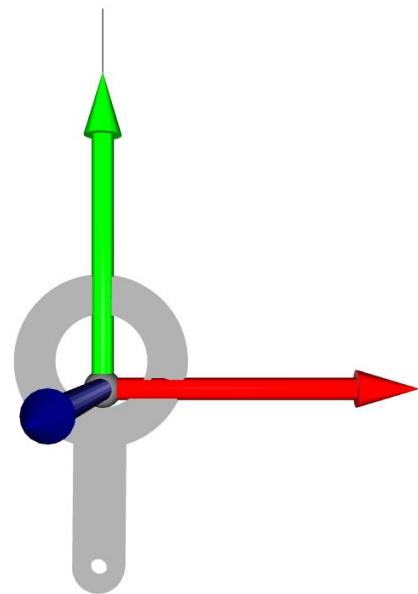
[2] <https://physics.stackexchange.com/questions/426273/how-to-find-the-axis-with-minimum-moment-of-inertia>

[3] <https://stackoverflow.com/questions/67017134/find-rotation-matrix-to-align-two-vectors>

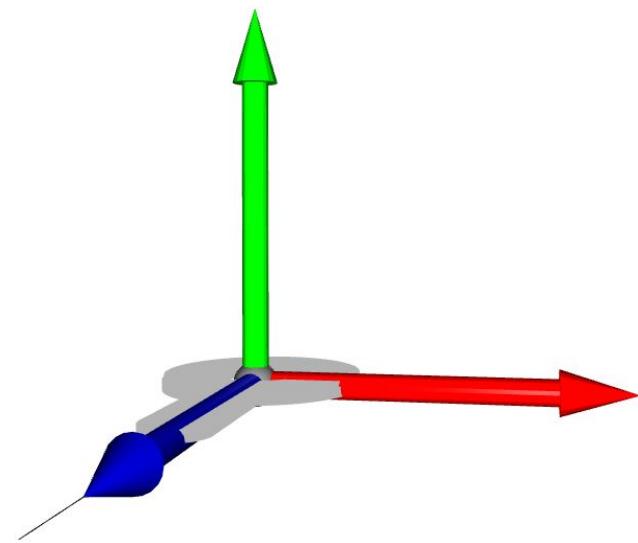
Alignment



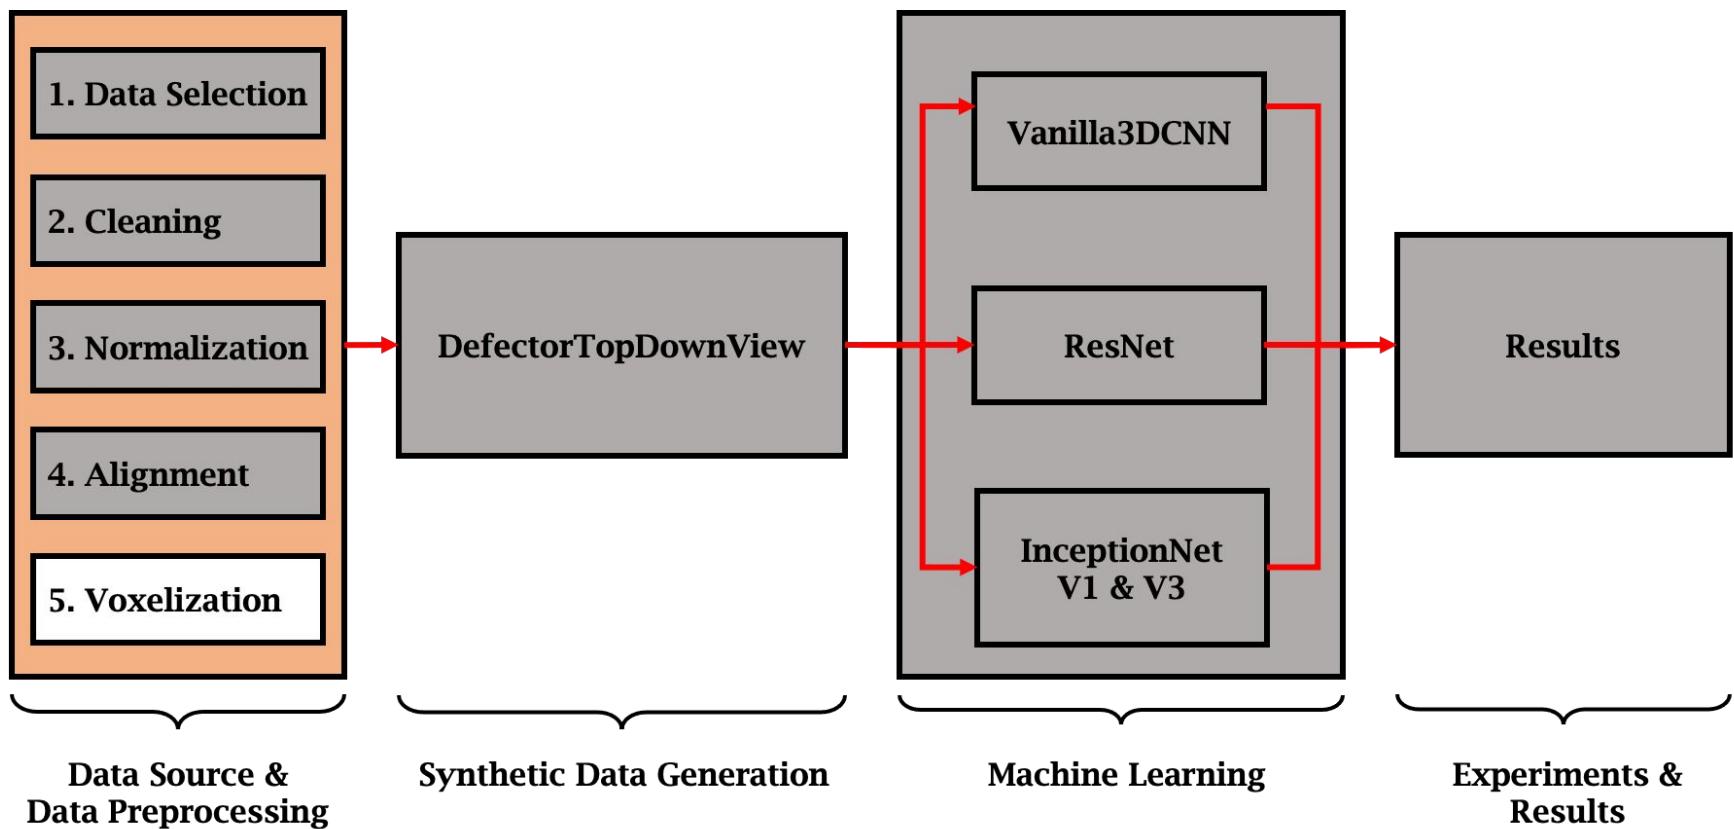
Alignment of min MOI axis and X-axis



Alignment of min MOI axis and Y-axis

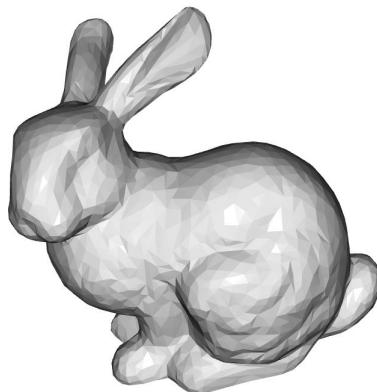


Alignment of min MOI axis and Z-axis

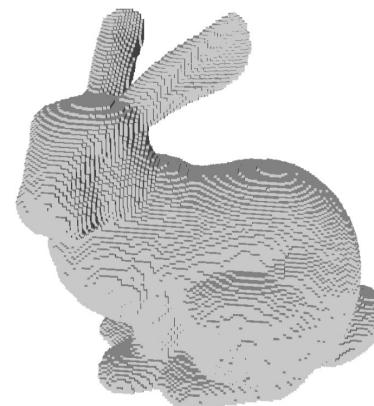


Voxelization

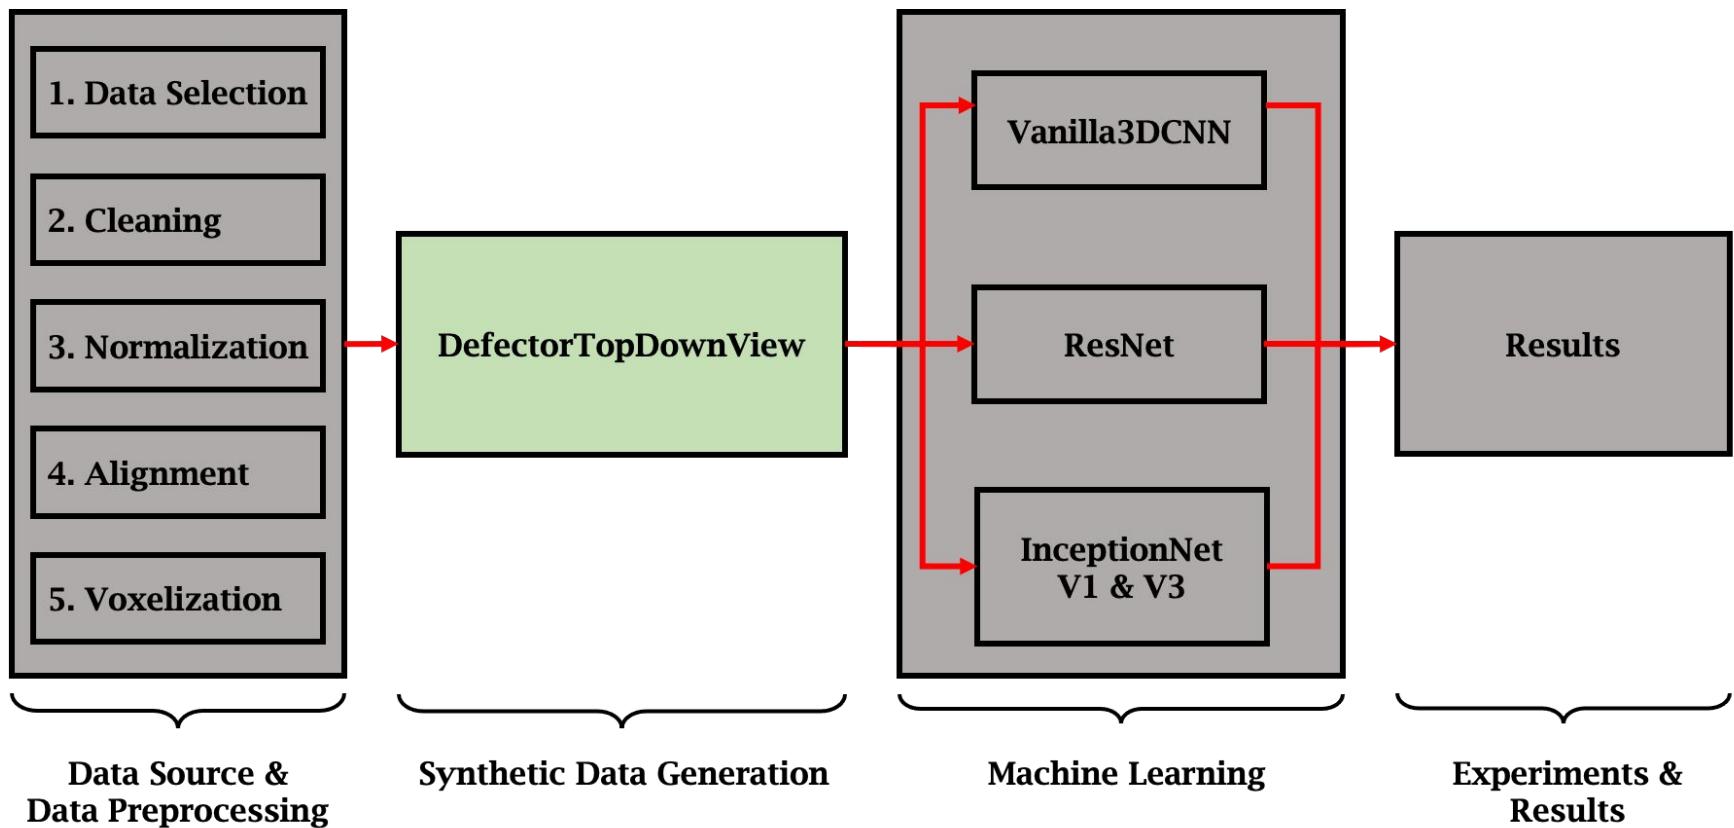
- All previous transformations were applied on the triangular mesh representation
- Final preprocessing step is converting the model to a voxelized representation
- What is a voxel representation?
 - A data representation that uses voxels. A voxel can be regarded as a pixel in a three-dimensional space.
- Why use voxel representation?
 - We hypothesized that SOTA CNNs will offer good performance after replacing 2D convolutions with 3D convolutions
 - It allows for an easy introduction of defects



Triangular Mesh Representation



Voxelized Representation

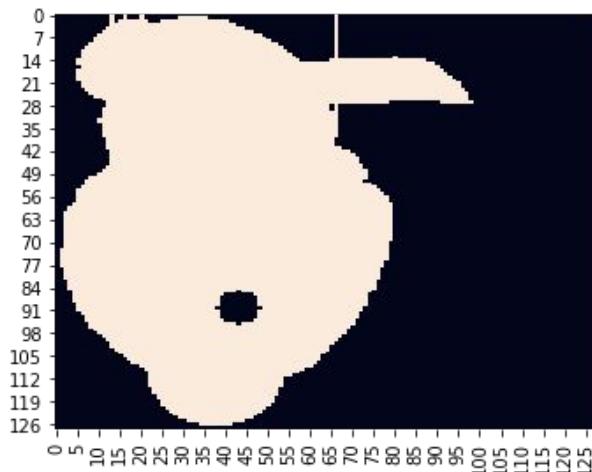


Assumptions & Definitions 1/2

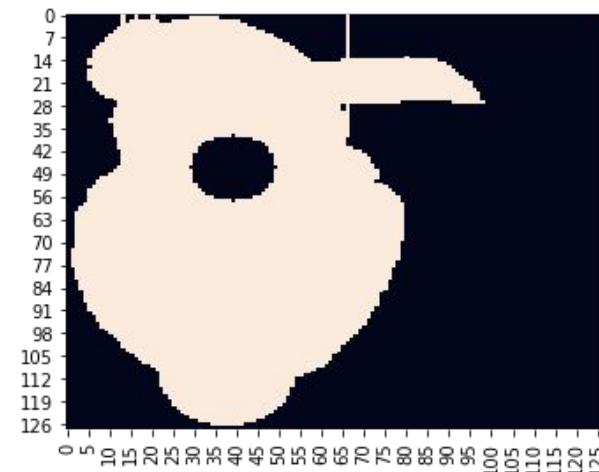
- **Problem:** Limited labeled data
- **Goal:** Synthetically add defects to 3D models in order to generate non-printable models
- **Assumption:**
 - All previous selected models are printable
 - Defects are here **holes aligned with the z-axis** that are added to the 3D model
 - Deciding parameters:
 - **Radius**
 - radius printable (10 voxels)
 - radius non-printable (5 voxels)
 - **Border**
 - Border printable (5 voxels)
 - Border non-printable (3 voxels)

Assumptions & Definitions 2/2

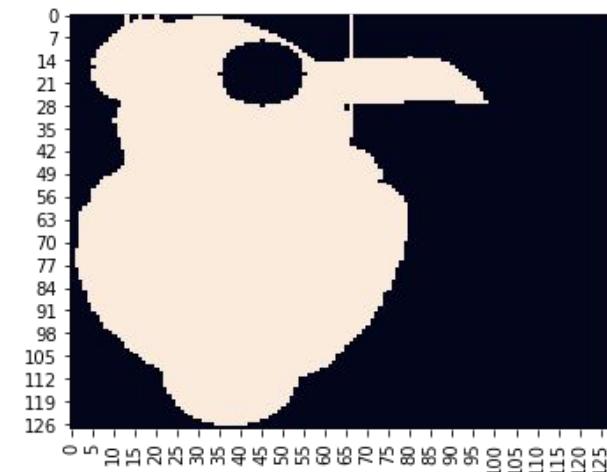
Resulting models:



Model **non-printable**
defect in middle



Model **printable** defect in
middle



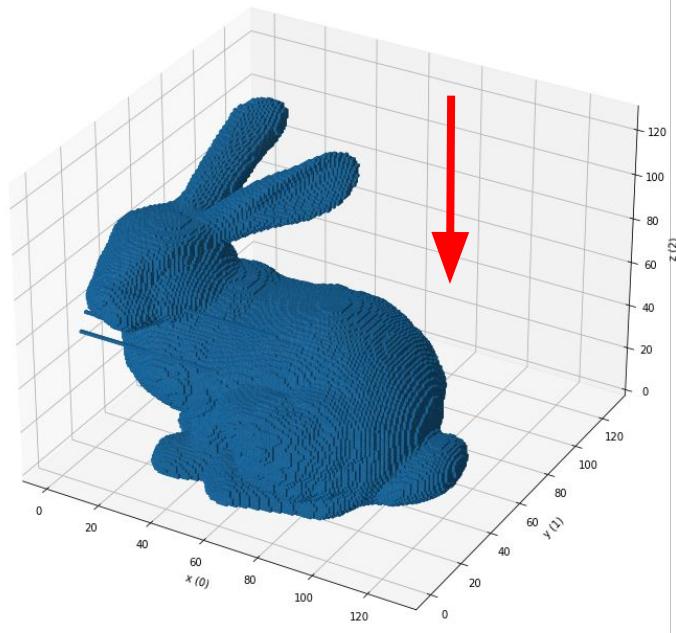
Model **non-printable**
defect at border

Idea

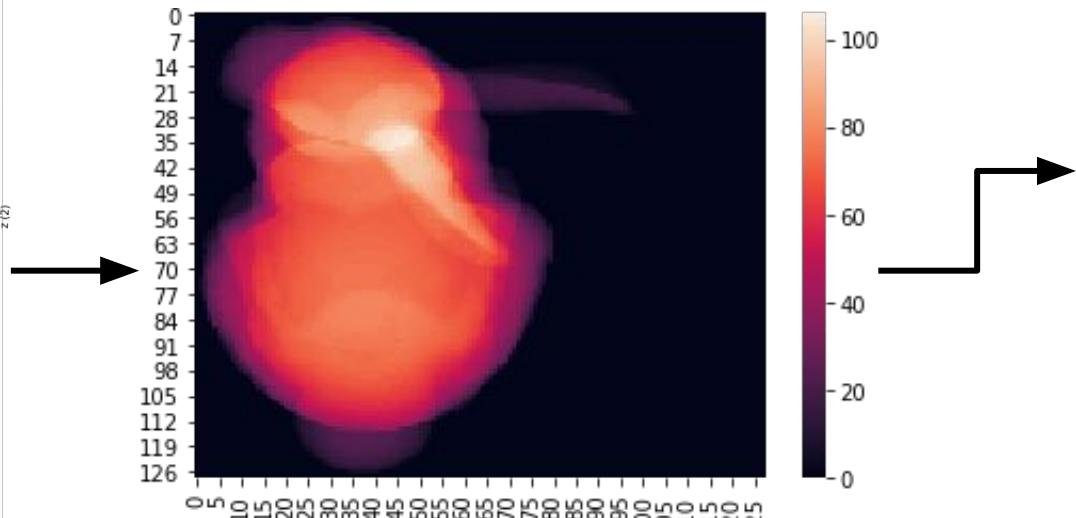
- Main Idea: Transform 3D model data s.t. it can be used to find right offset

→ TopDownView:

- Inspiration: **Heatmaps** (visualizing 3D data in 2D), **common sense**
- **Project 3D model data onto the (x, y)-grid**

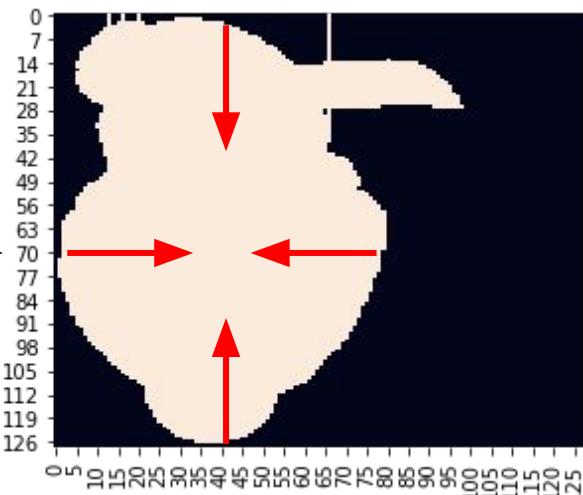


Summation over the z-axis

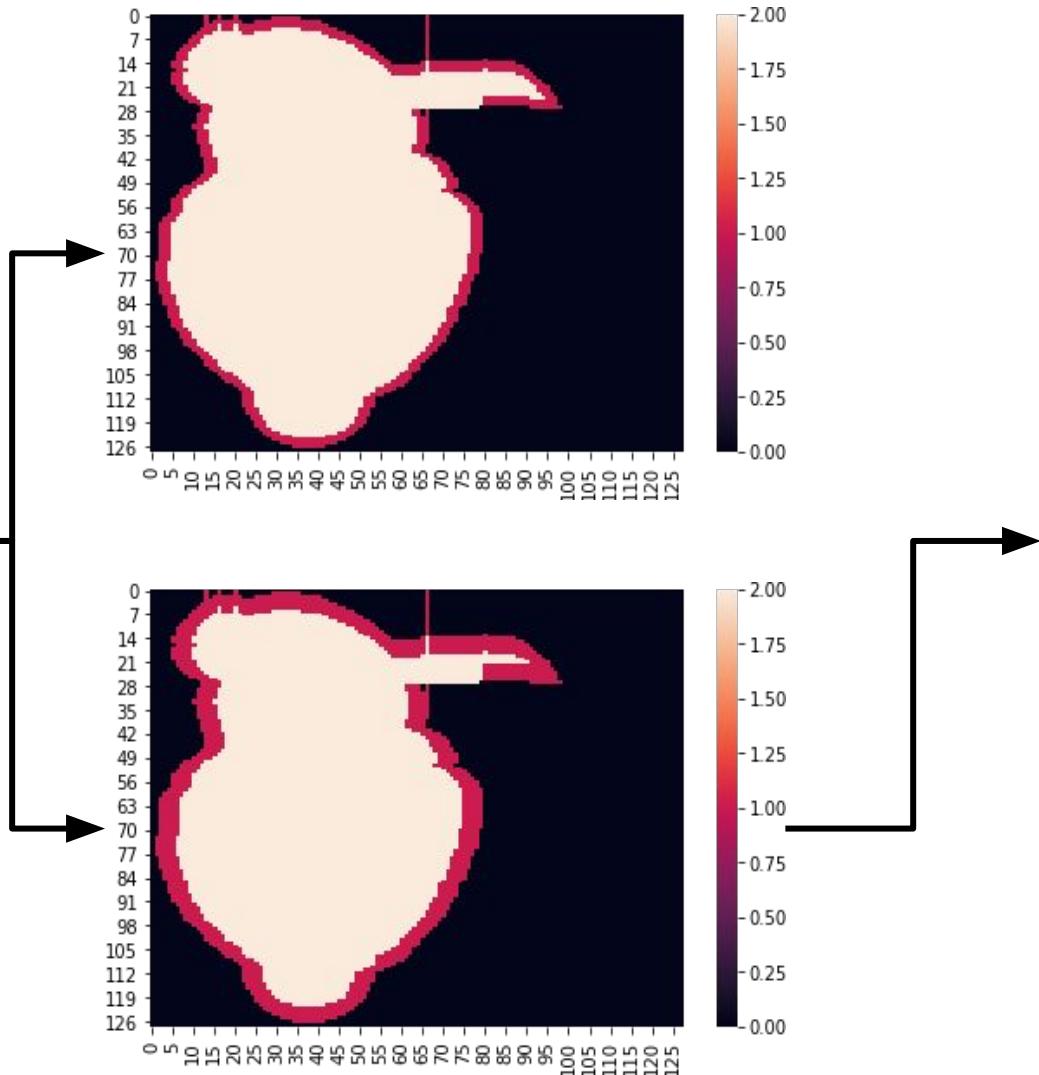


TopDownView

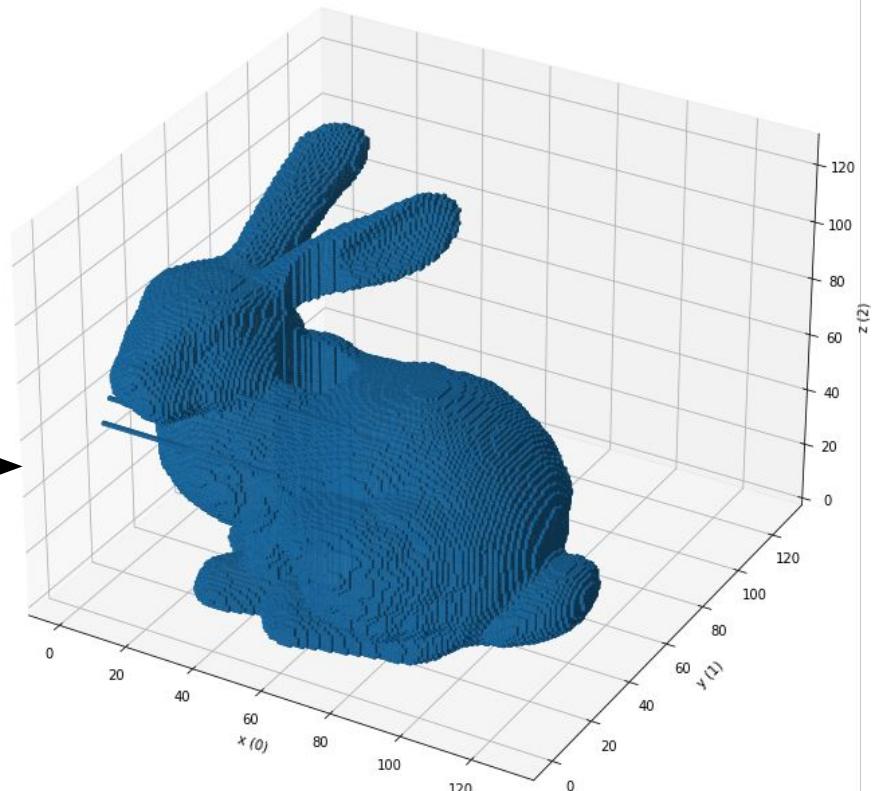
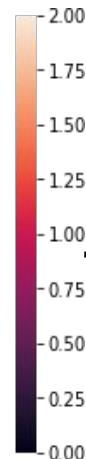
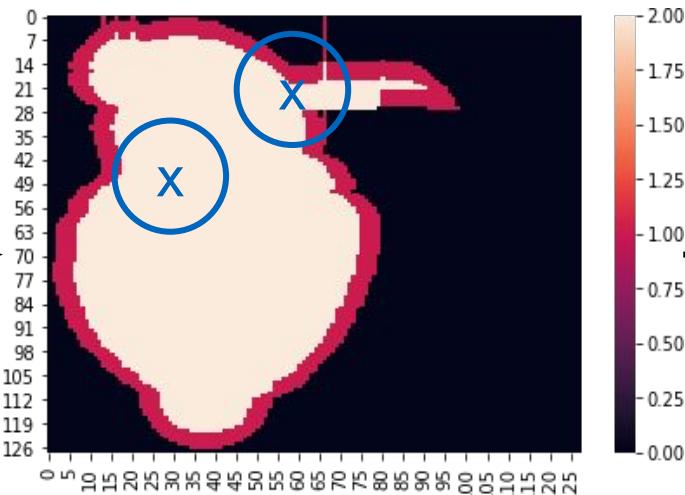
Defects in the middle 1/2



Offset Preselection (X and Y direction)

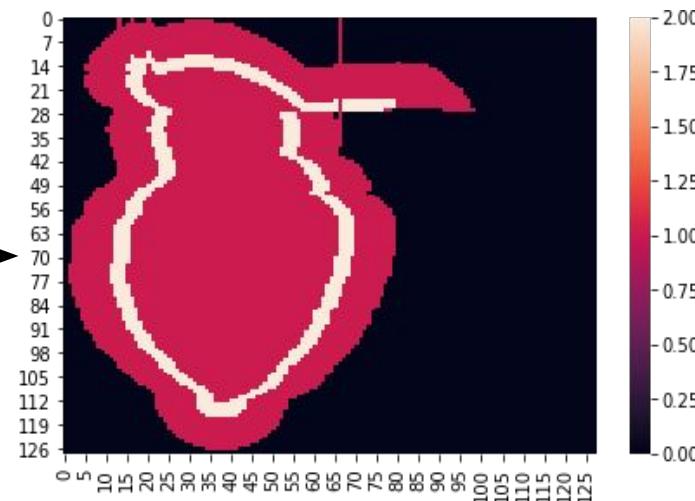
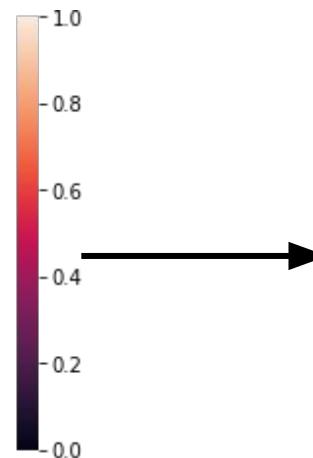
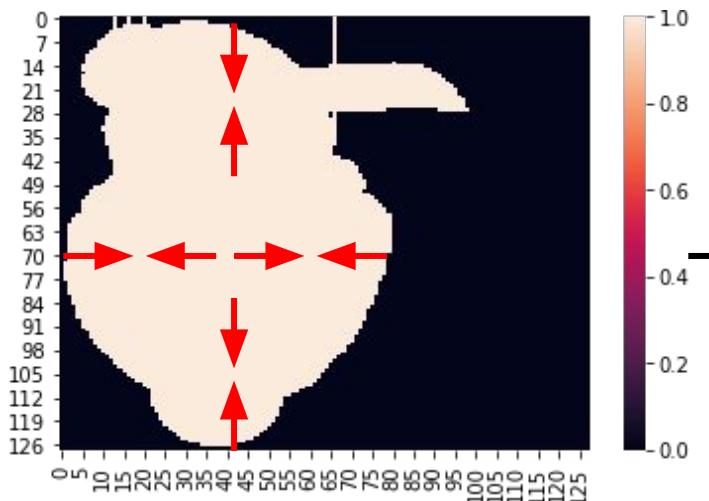


Defects in the middle 2/2



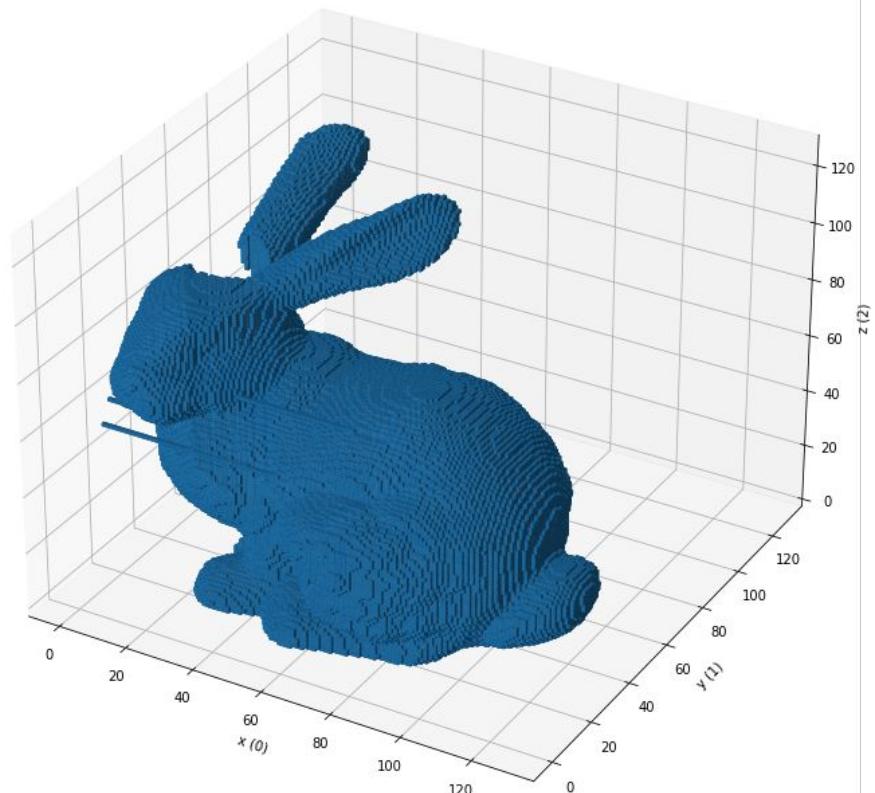
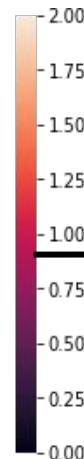
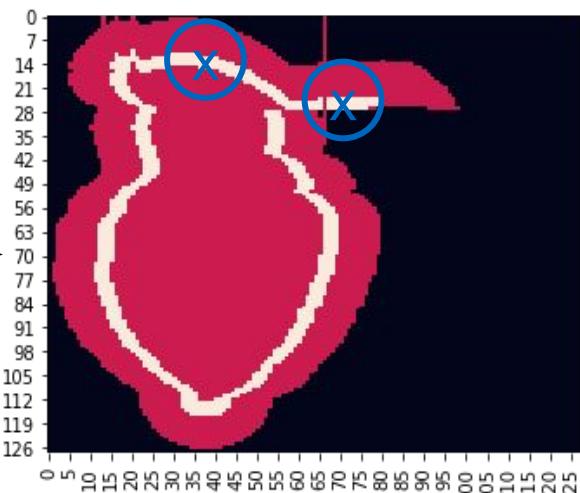
Non-axis direction check

Defects at the border 1/2



Offset Preselection (X and Y direction)

Defects at the border 2/2

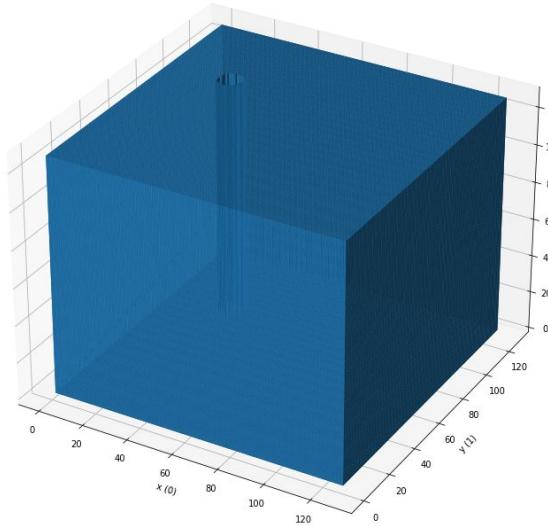


Non-axis direction check

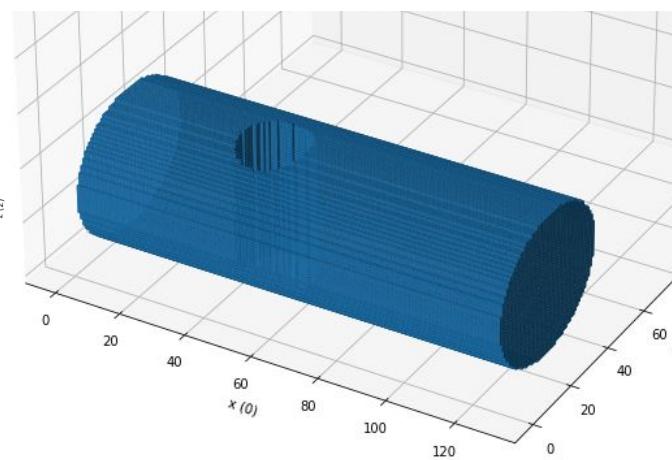
AMC Dataset

→ Additive Manufacturing Classification (AMC) Dataset:

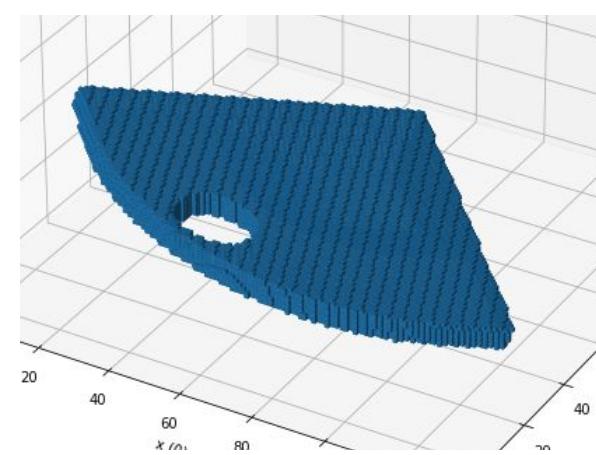
- 7430 3D models
- Balanced in terms of labels (printable / non-printable)
- But: Defects not equally distributed



Non-printable defect middle



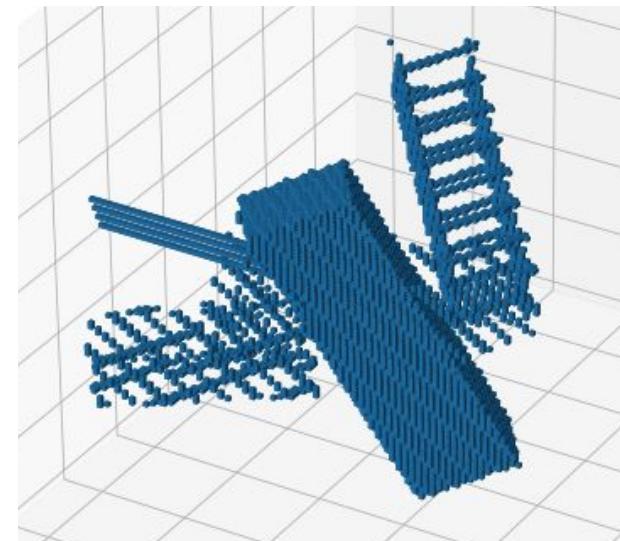
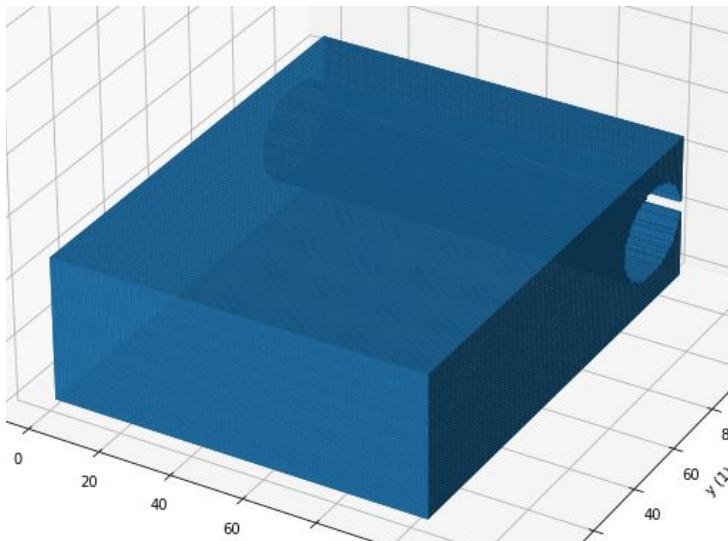
Printable defect middle

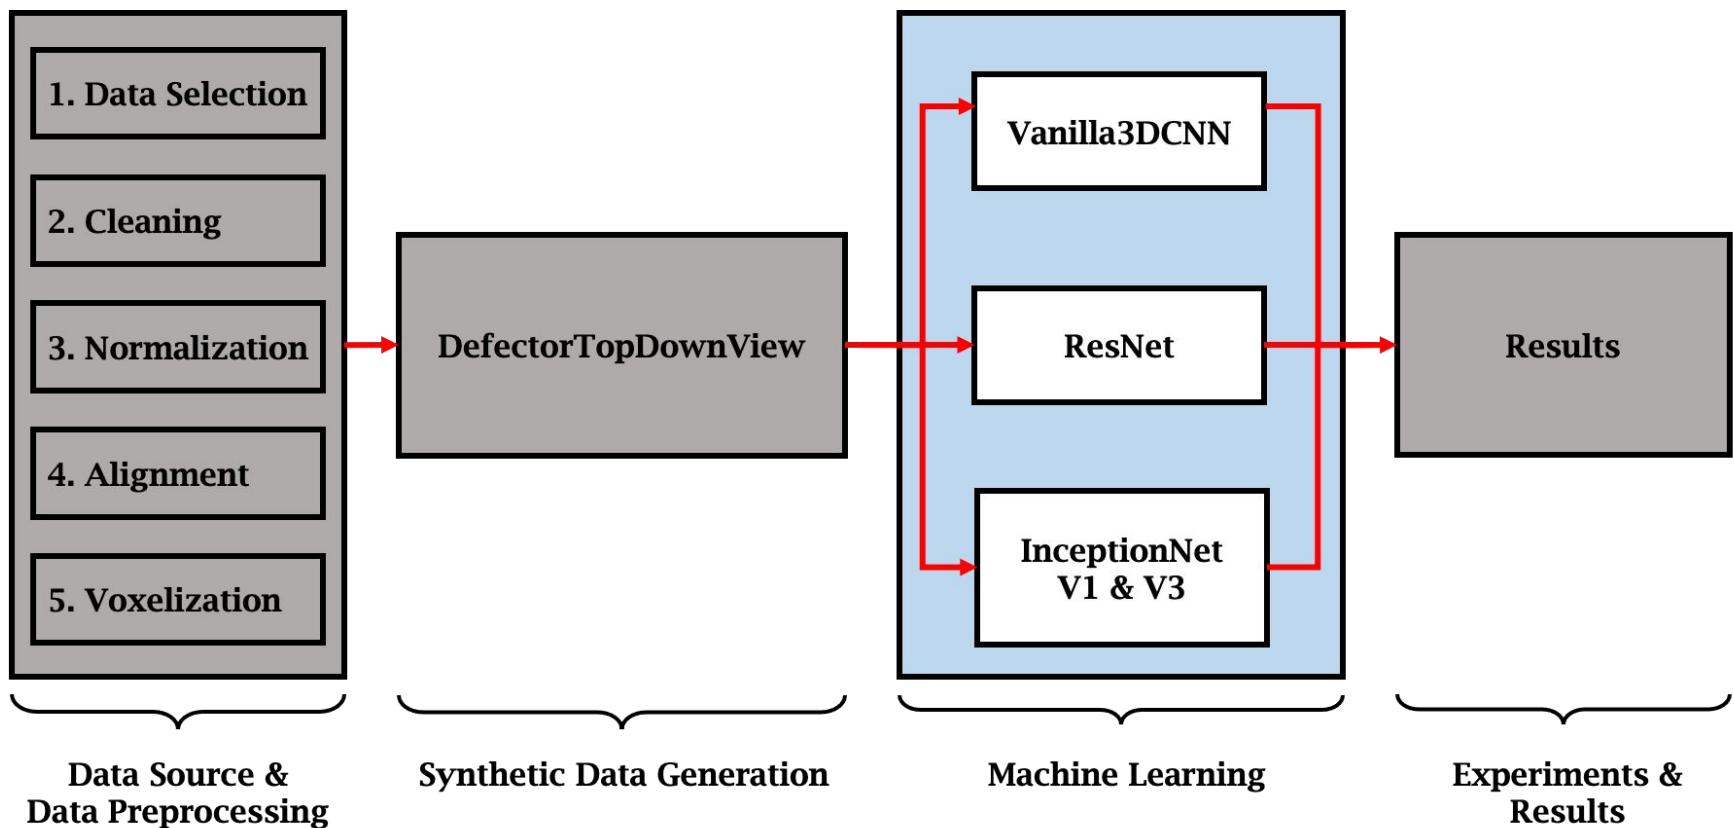


Non-printable defect border

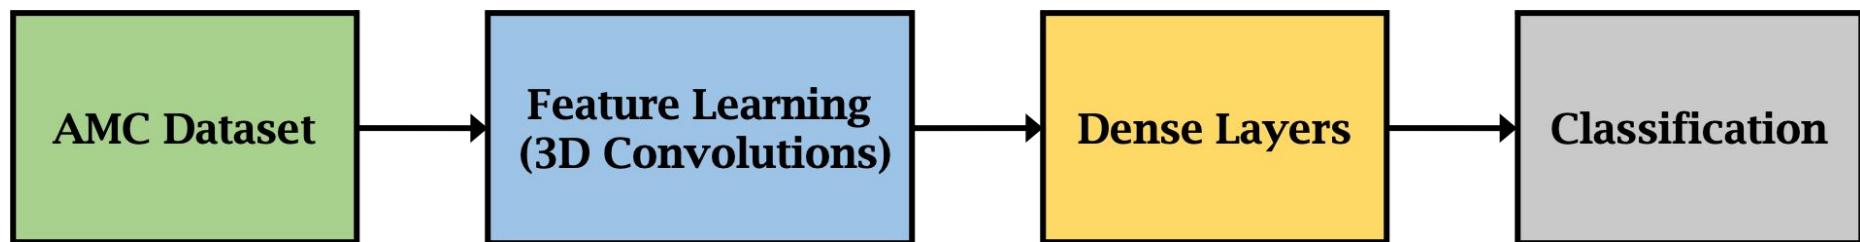
Limitations

- **Strong assumptions** on the input models
 - **Selected models already contain holes or are too complex**
- **Artificial defined defects**, i.e. too far away from the real given problem, strong abstraction
- If models and the defects gets more complex, this approach **quickly reaches its limitations**

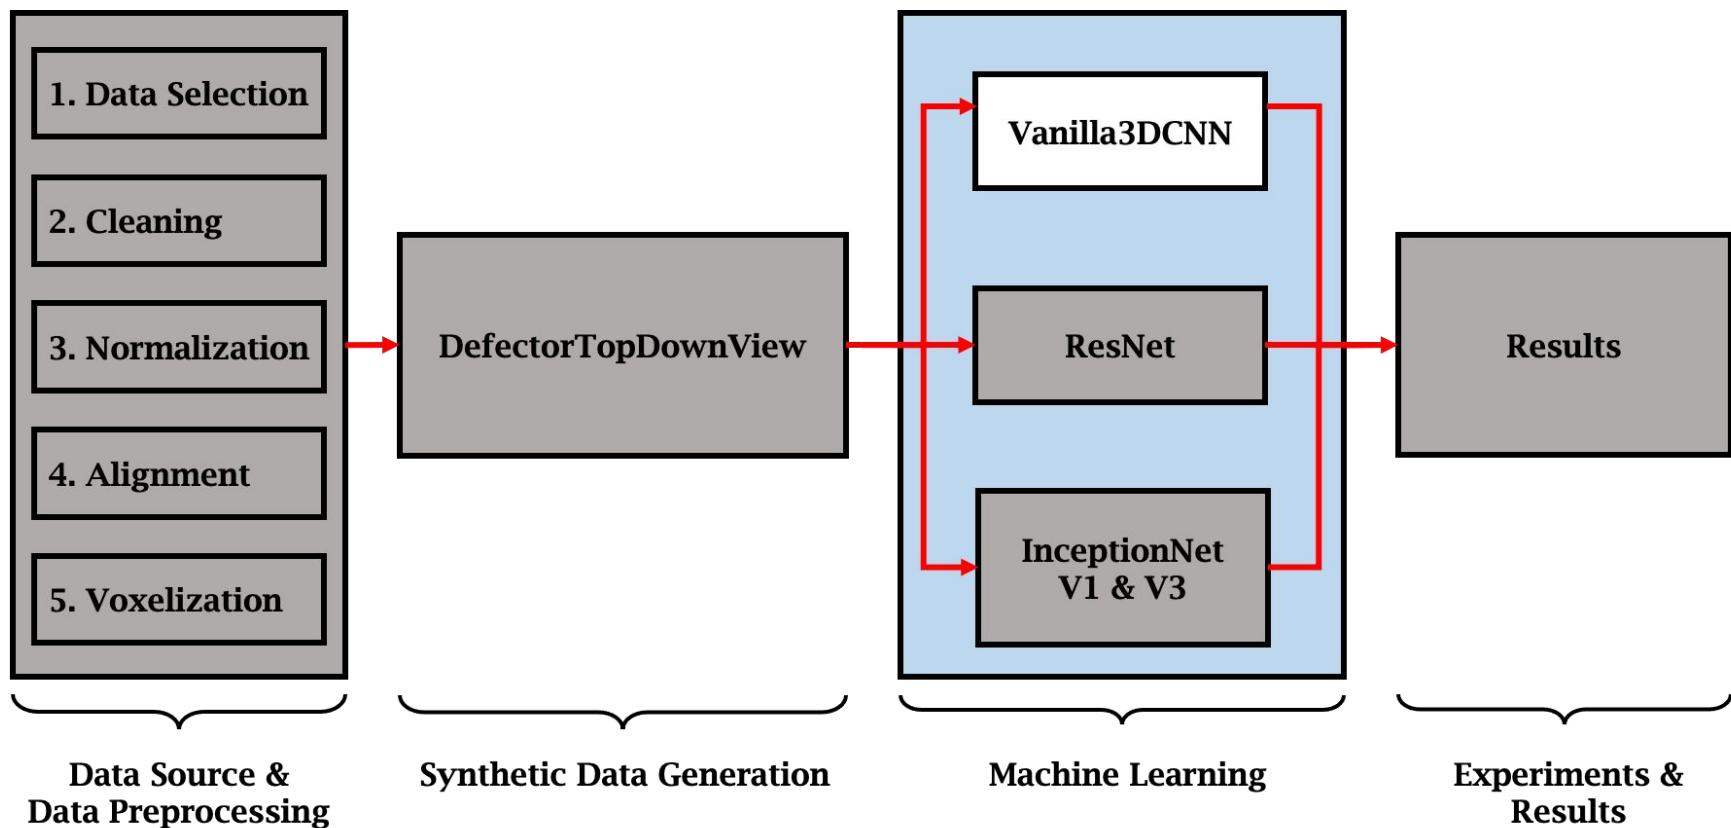




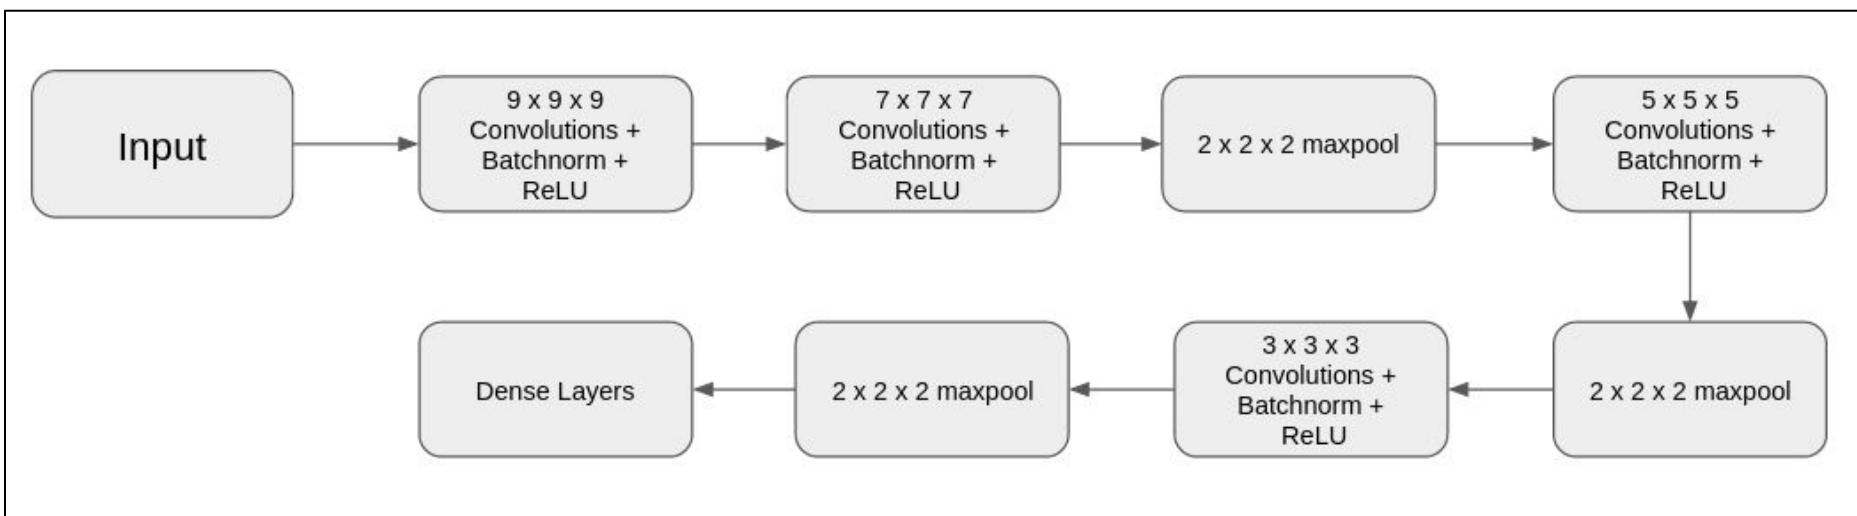
Deep Learning Workflow



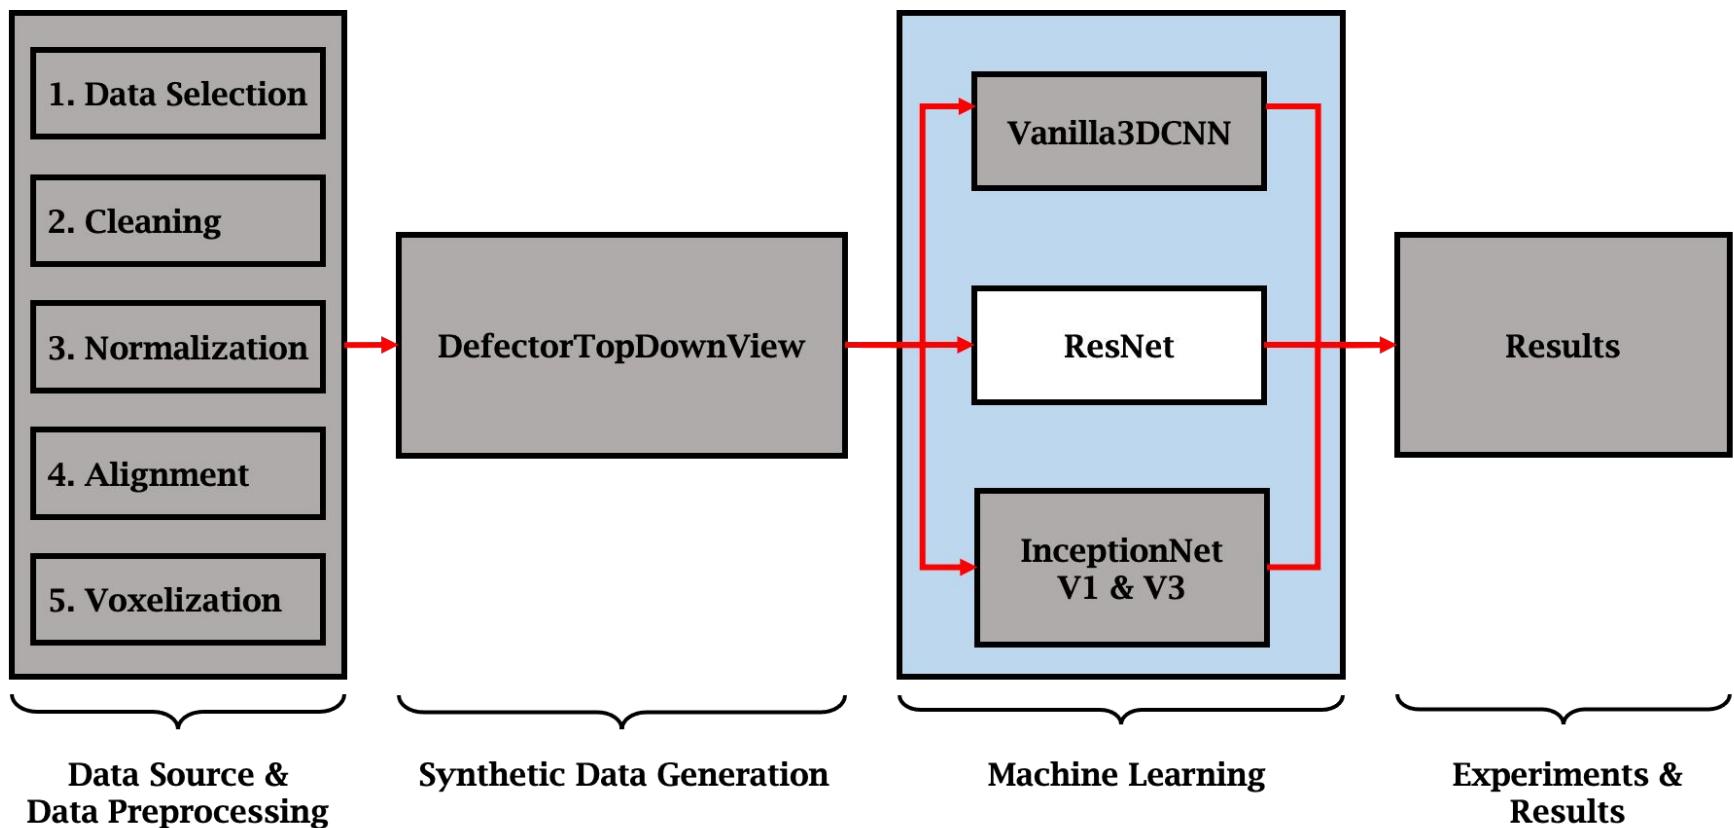
- Type of problem to be solved: Binary Classification
- Architectures examined for feature extraction:
 - Vanilla3DCNN
 - ResNet
 - InceptionNet V1
 - InceptionNet V3



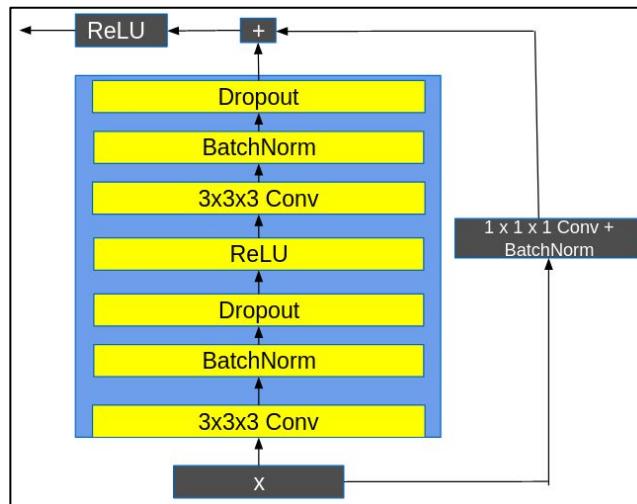
Vanilla3DCNN Architecture



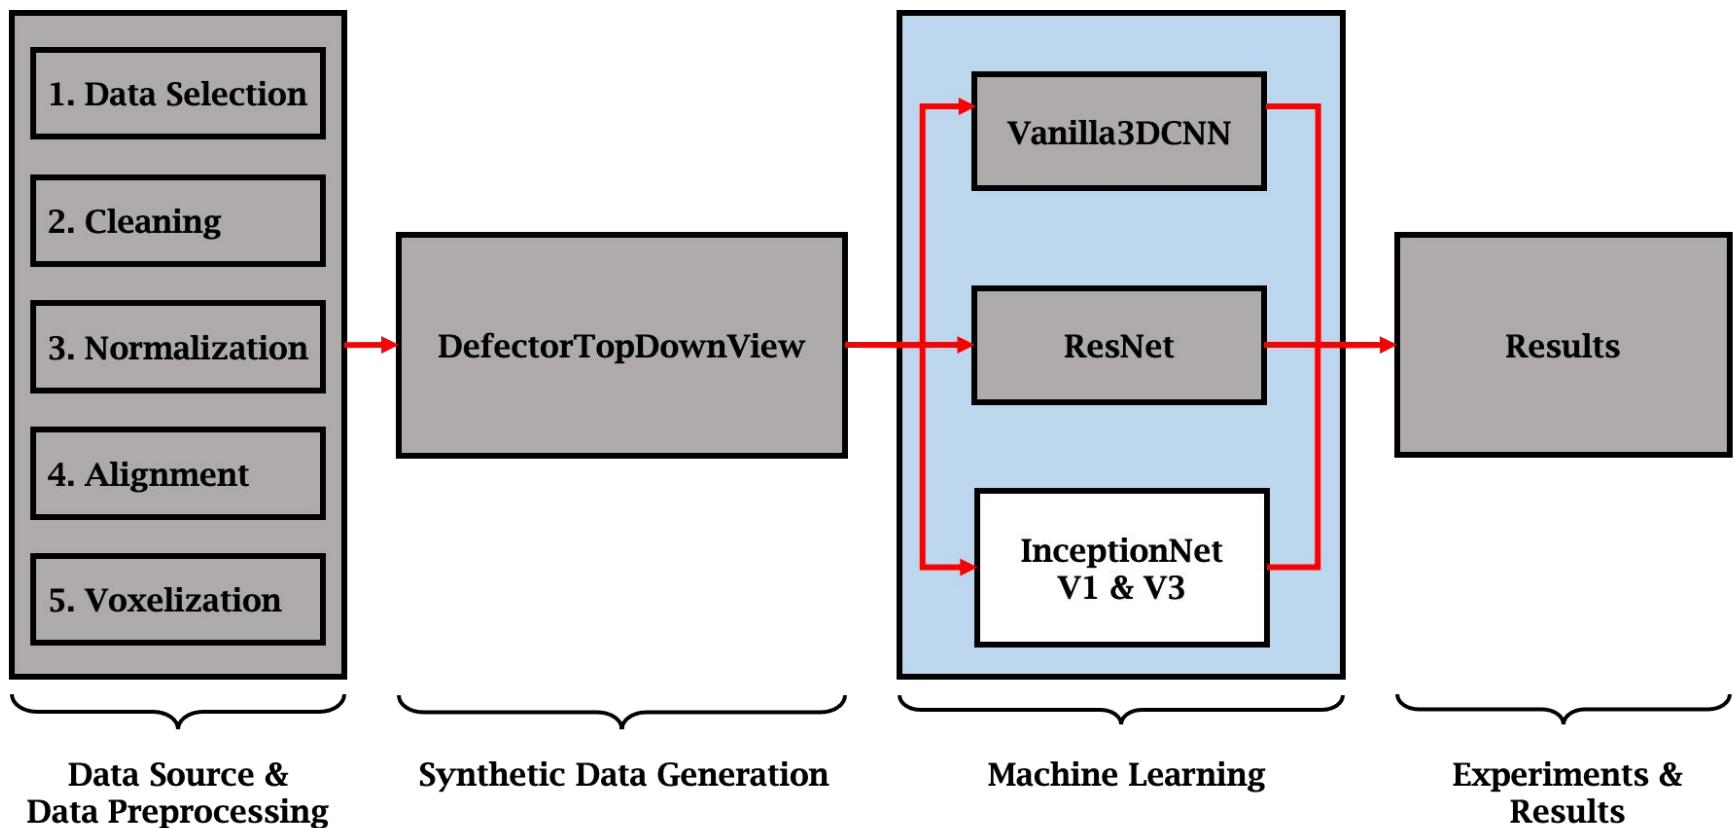
- Number of trainable parameters: **1.8M**
- 4 convolution layers with kernel sizes 9, 7, 5 and 3 were used.
- Maxpooling layer to **reduce spatial size**.
- Batch normalization layers to **stabilize the training process**.



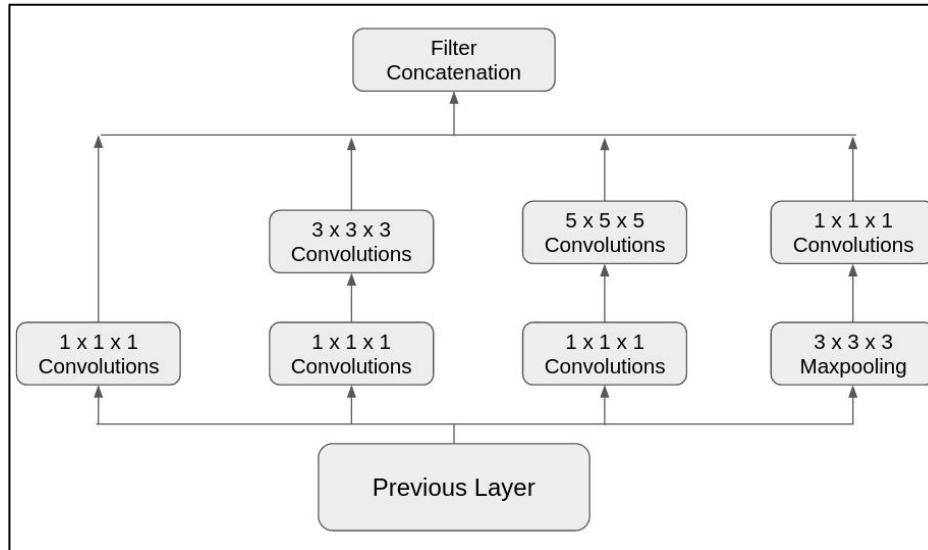
ResNet Architecture



- Number of trainable parameters: **9.1M**
- Basic residual block contains two $3 \times 3 \times 3$ layers with batch norm layers and dropout layers.
- Basic residual block repeated twice.
- To **avoid sudden reduction of input spatial dimensions**, two pooling layers used at the end.

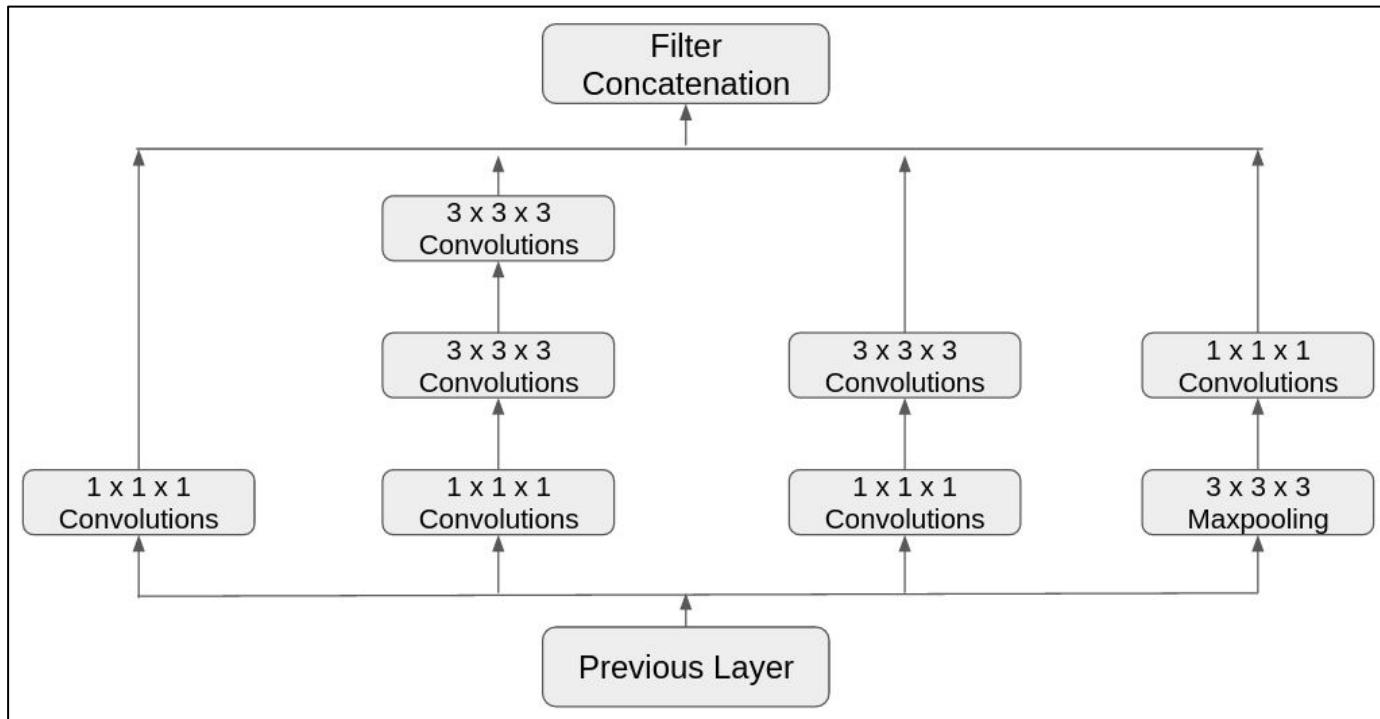


InceptionNet V1 Architecture



- Basic intuition: **Let's go wider.**
- Number of trainable parameters: **9.6M**
- To capture the salient features having varying dimensions, choosing right kernel size is difficult.
- Each InceptionNet V1 block has kernel sizes 1,3 and 5 operating at the same level.
- InceptionNet V1 block repeated 4 times.

InceptionNet V3 Architecture



- Number of trainable parameters: **17.9M**
- Improvement over InceptionNet V1
- Disadvantage of InceptionNet V1: **Large reduction** in spatial dimension due to kernel size $5 \times 5 \times 5$.

Implementation Details

1. Activation Function:

- ReLU^[1] activation for all layers except for the final layer.
- Sigmoid^[1] activation for the final layer.

2. Loss Function:

- Binary cross entropy loss^[2] used.

$$BCE = - \sum_{i=1}^{C=2} t_i \log(f(s_i)) = -t_1 \log(f(s_1)) - (1 - t_1) \log(1 - f(s_1))$$

[1] Activation Functions: Comparison of Trends in Practice and Research for Deep Learning

[2] Generalized Cross Entropy Loss for Training Deep Neural Networks with Noisy Labels

Implementation Details

3. Weight Initialization:

- Kaiming normal^[3] weight initialization with **Fan-out mode** was used.
- Weights follow normal distribution as shown below:

$$W \sim N\left(0, \sqrt{\frac{2}{n}}\right)$$

4. Optimizer:

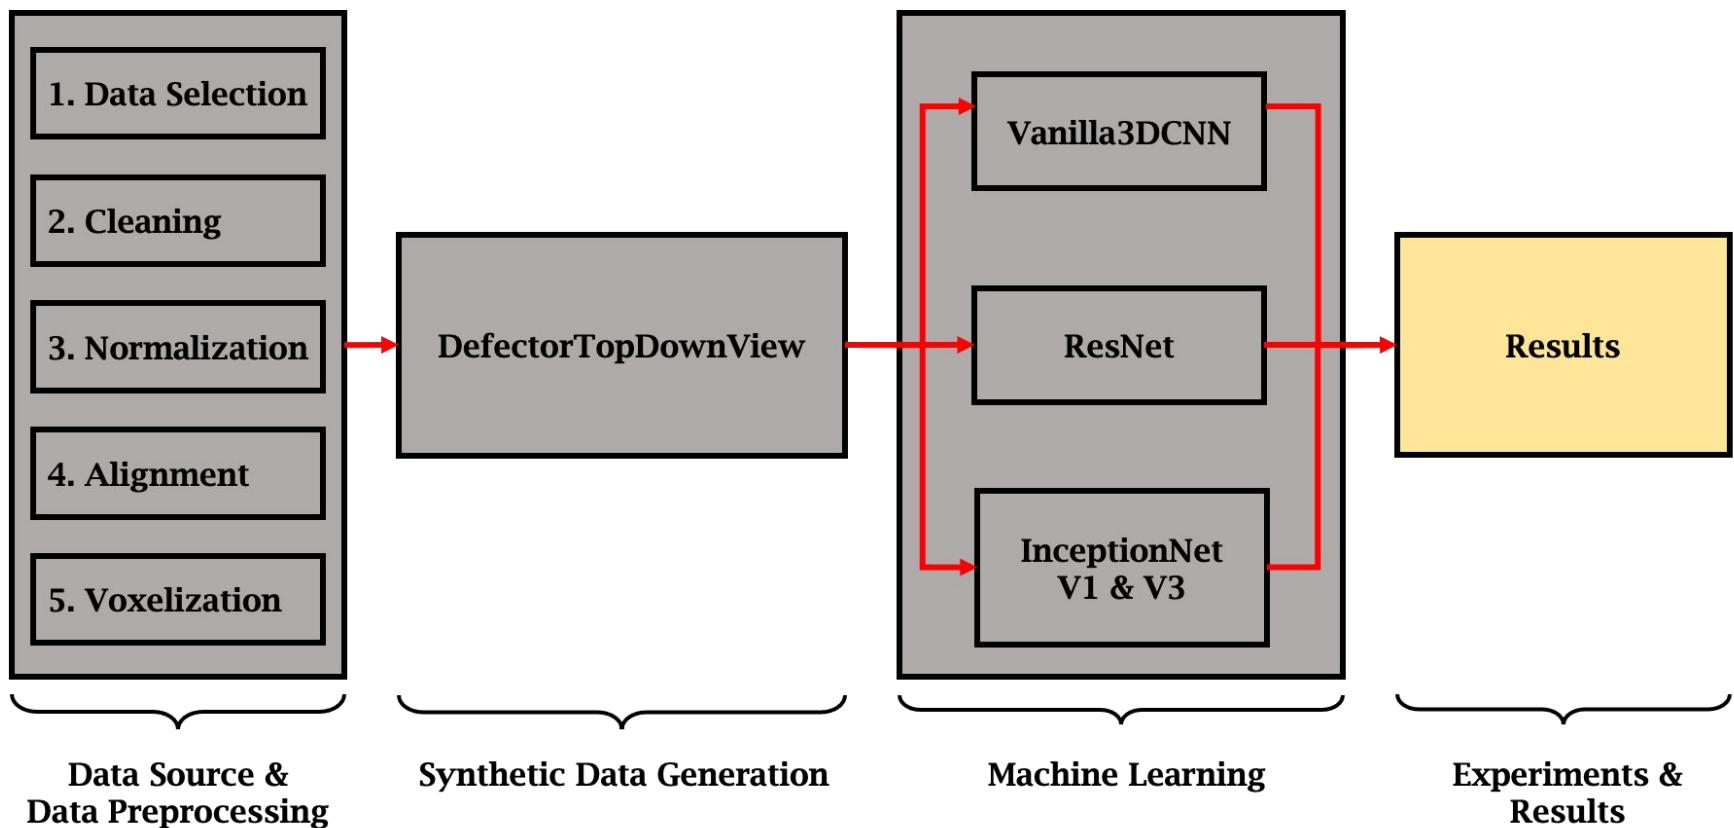
- Momentum based optimization algorithm is the default algorithm for all architectures.
- Adam optimizer^[4] used: Computes **adaptive learning rates** for each parameter.

4. Dataset:

- Total data samples: 7430
- Train/Validation split (random): 80/20, 5944 (train) and 1486 (validation)
- Law of large numbers: train and validation set are **balanced**

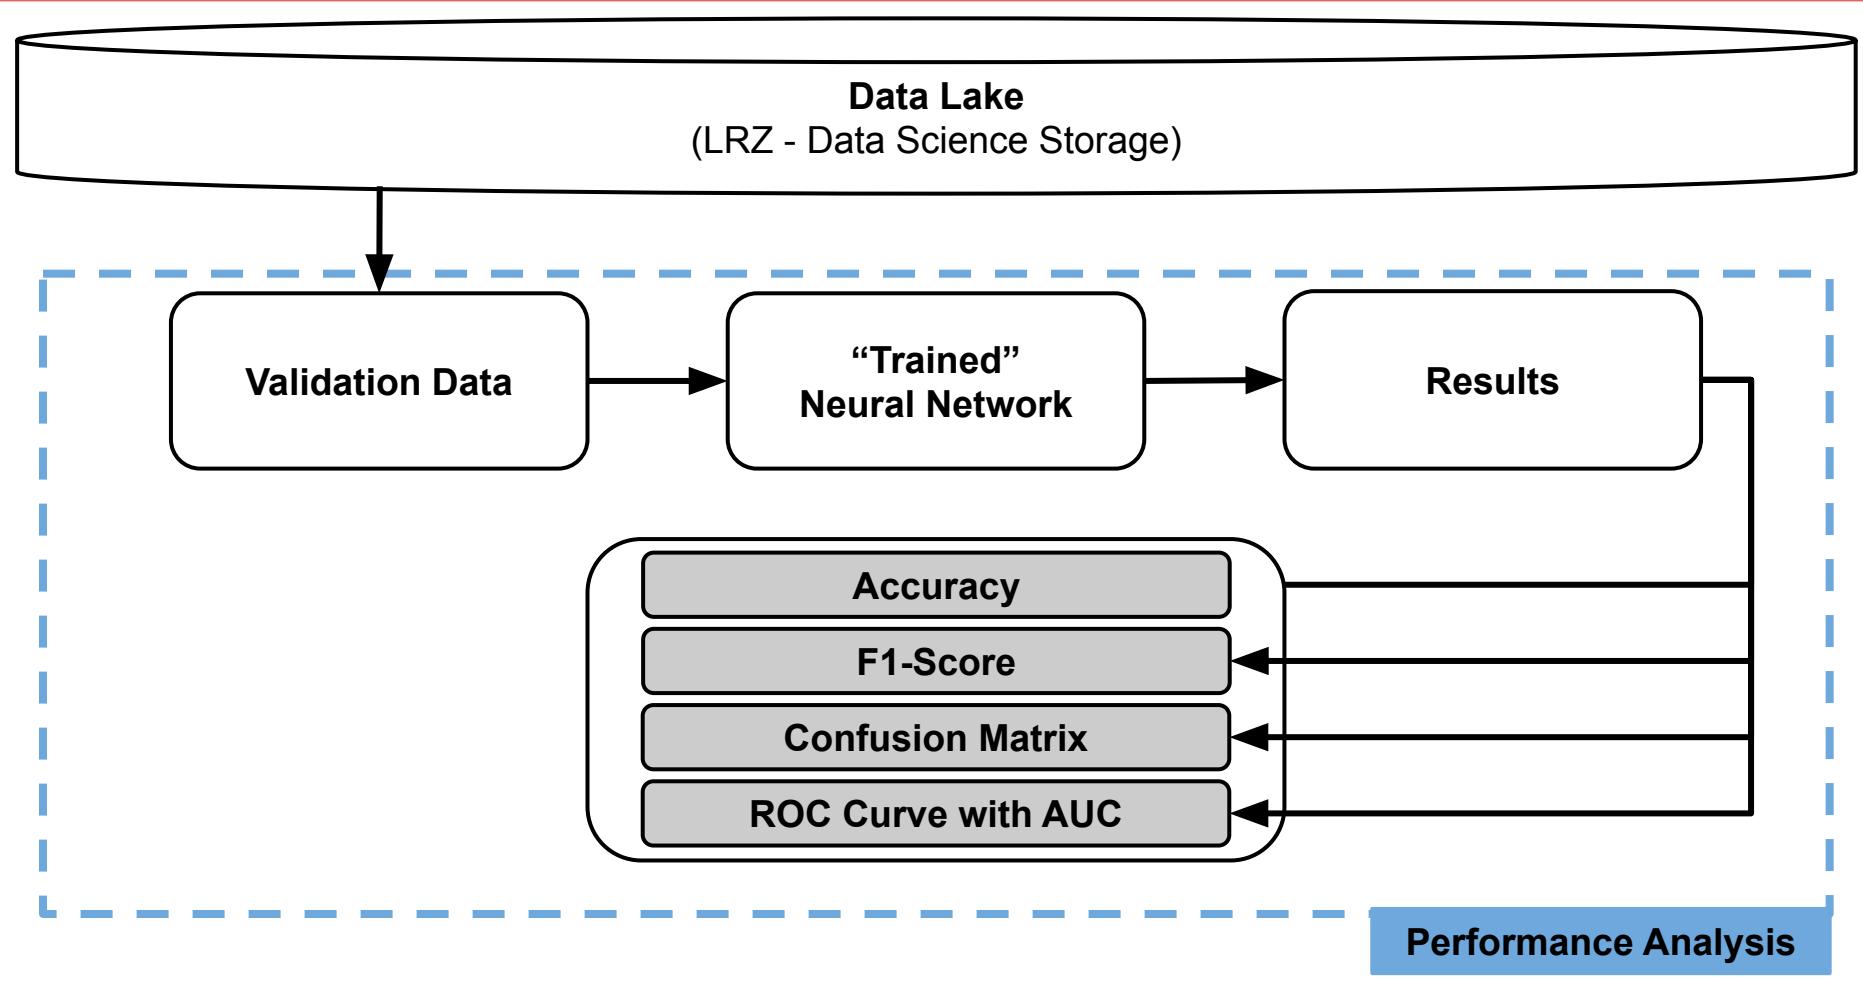
[3] Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification.

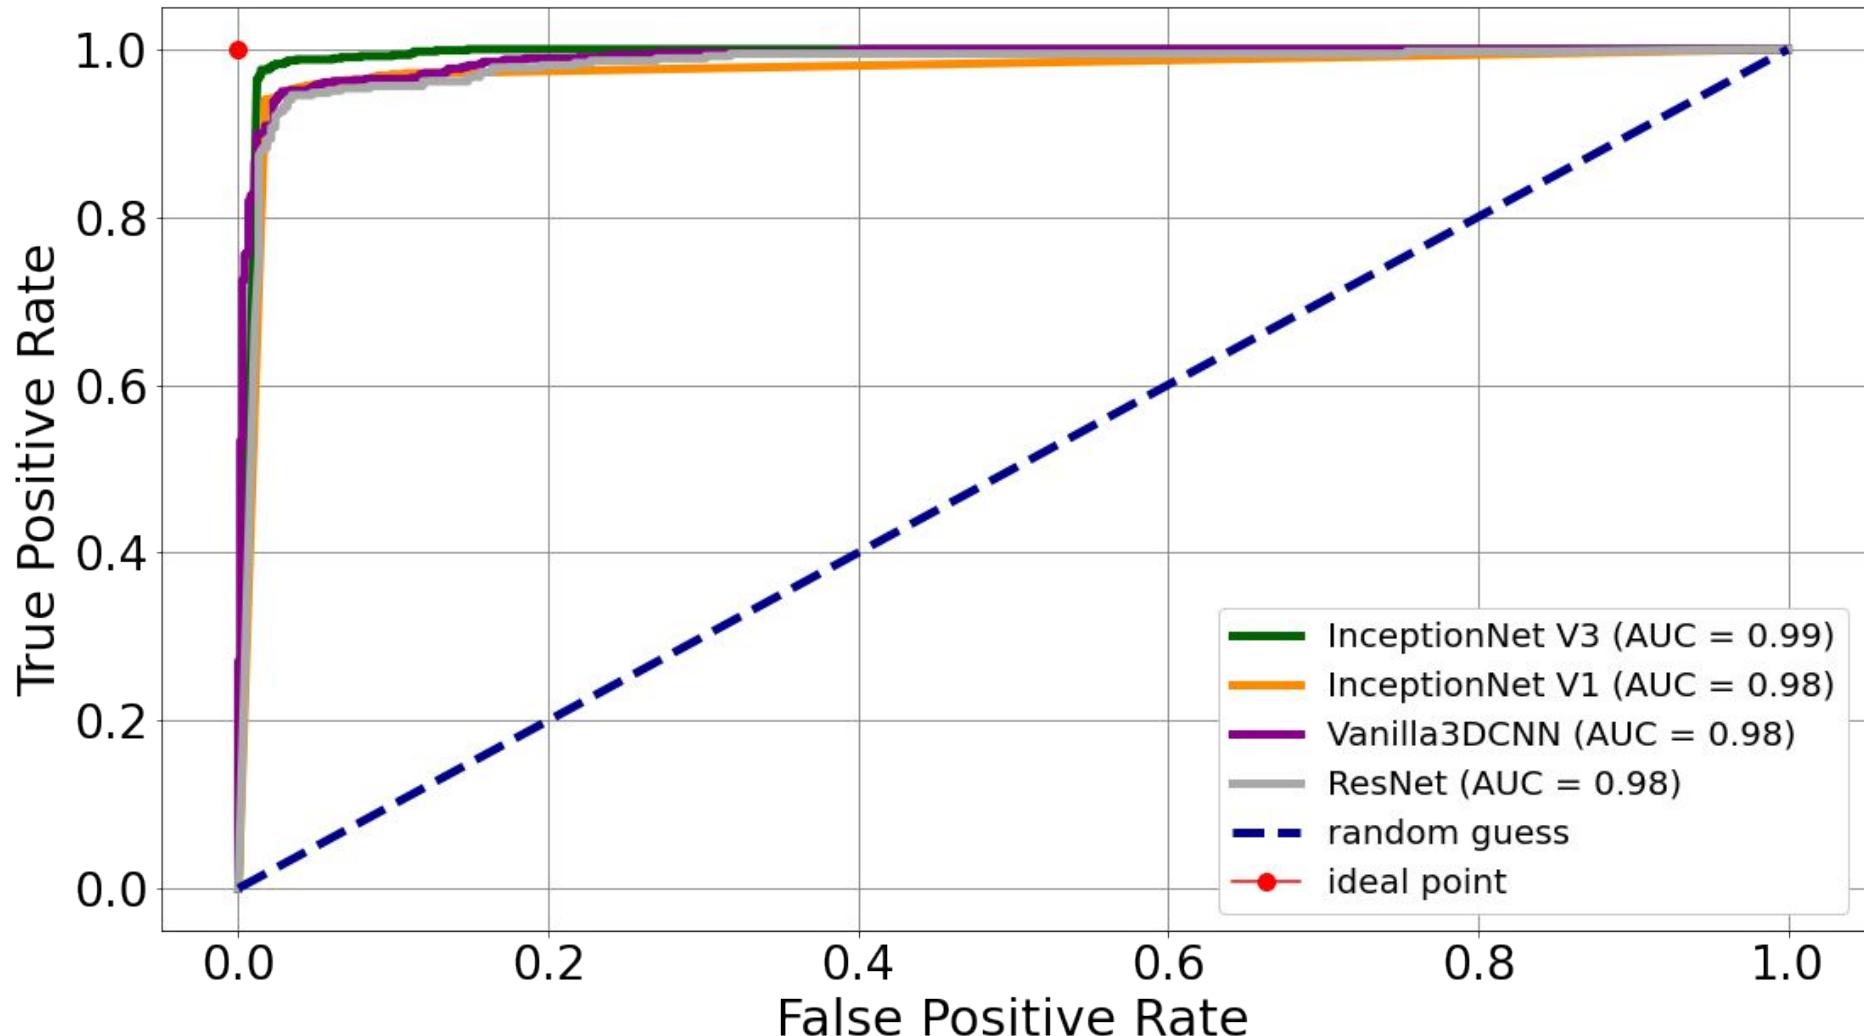
[4] Adam : A Method for Stochastic Optimization



Results & Performance Analysis

LRZ AI System

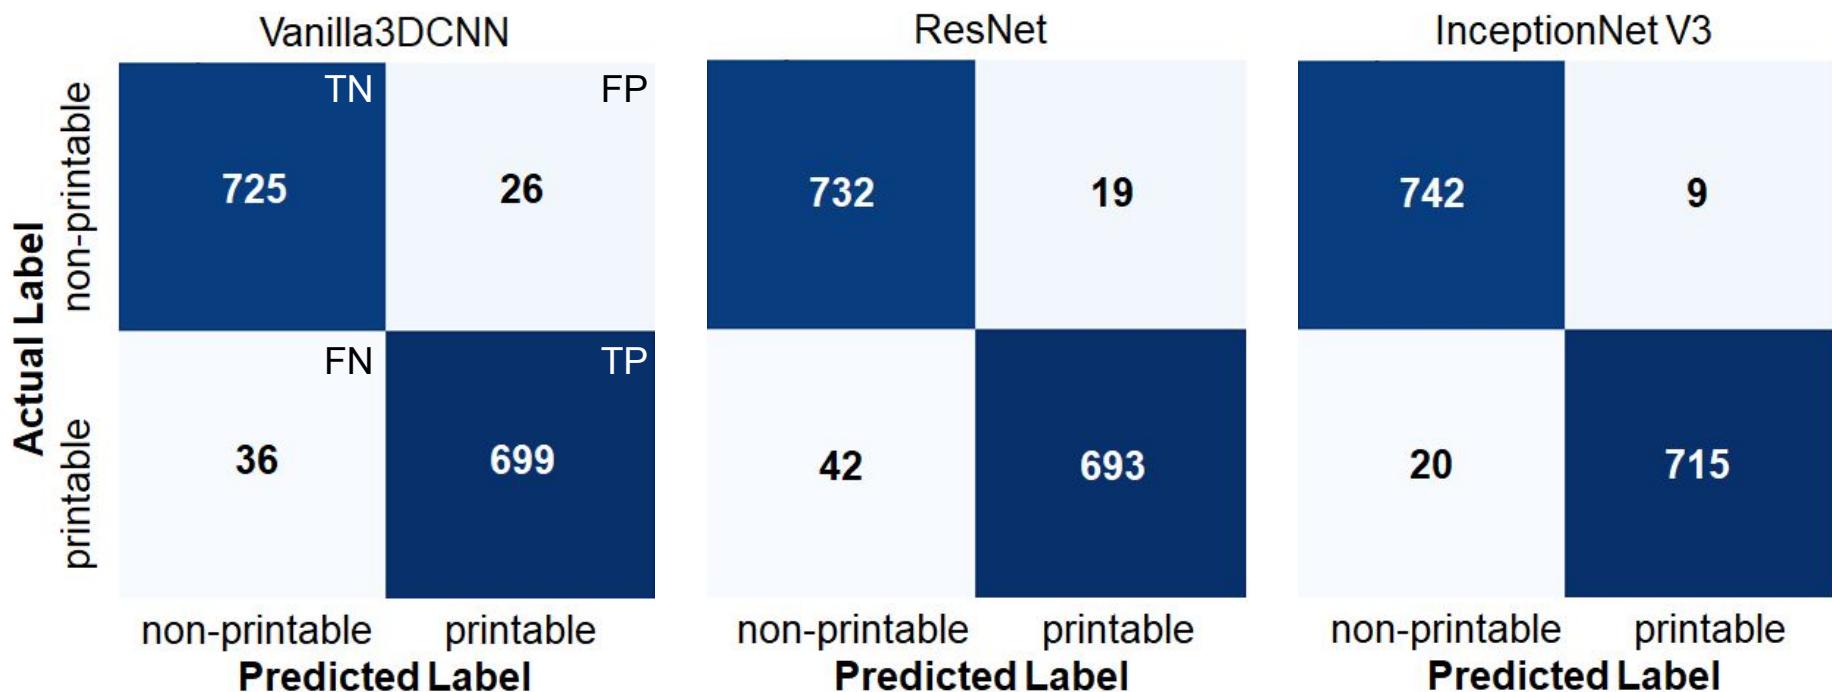


ROC⁽¹⁾ Curve with AUC⁽²⁾

(1) ROC: Receiver Operating Characteristic

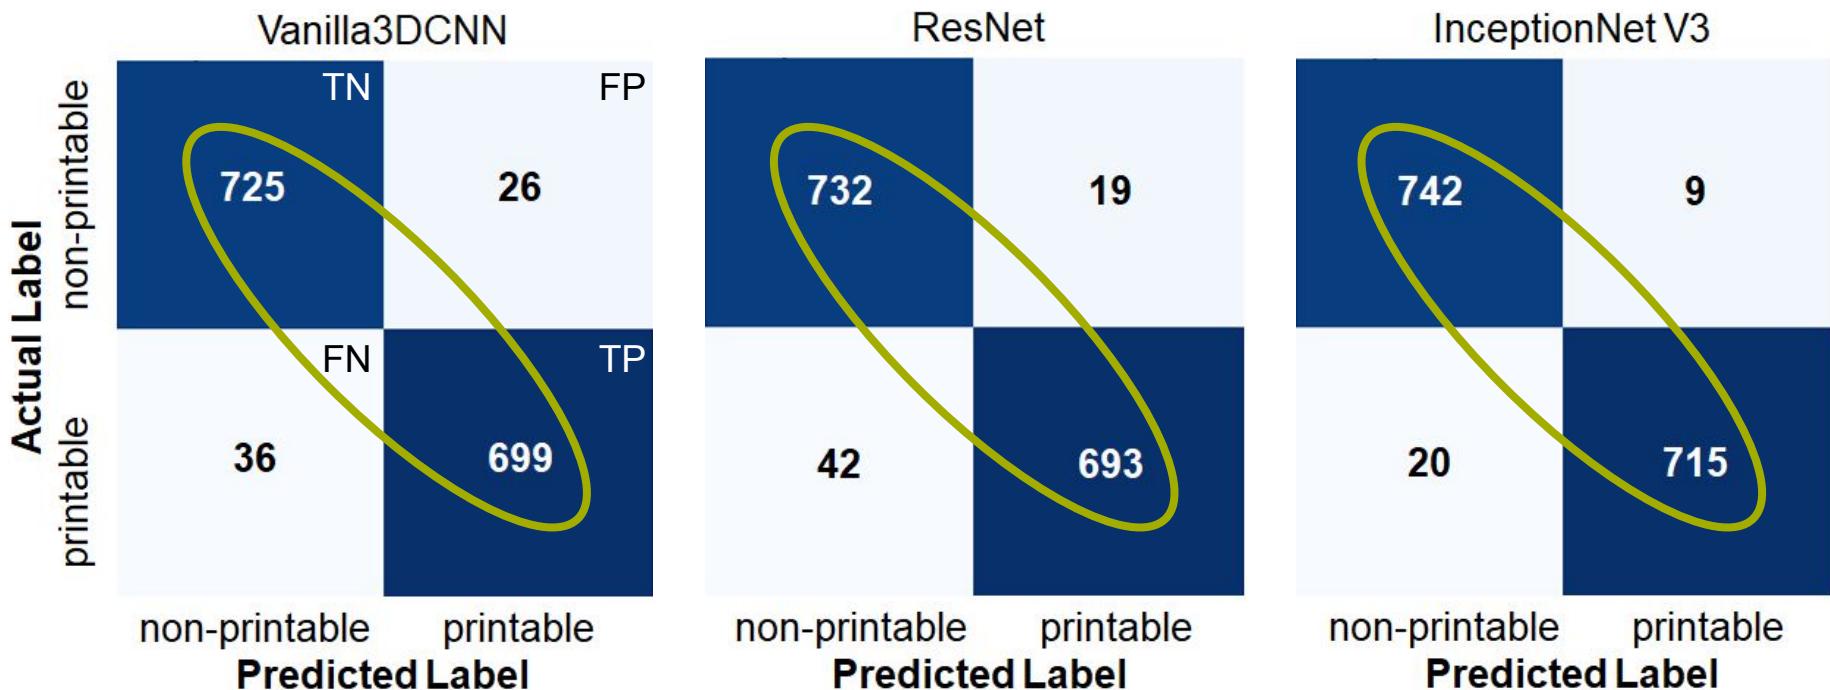
(2) AUC: Area under the ROC Curve

Confusion Matrix



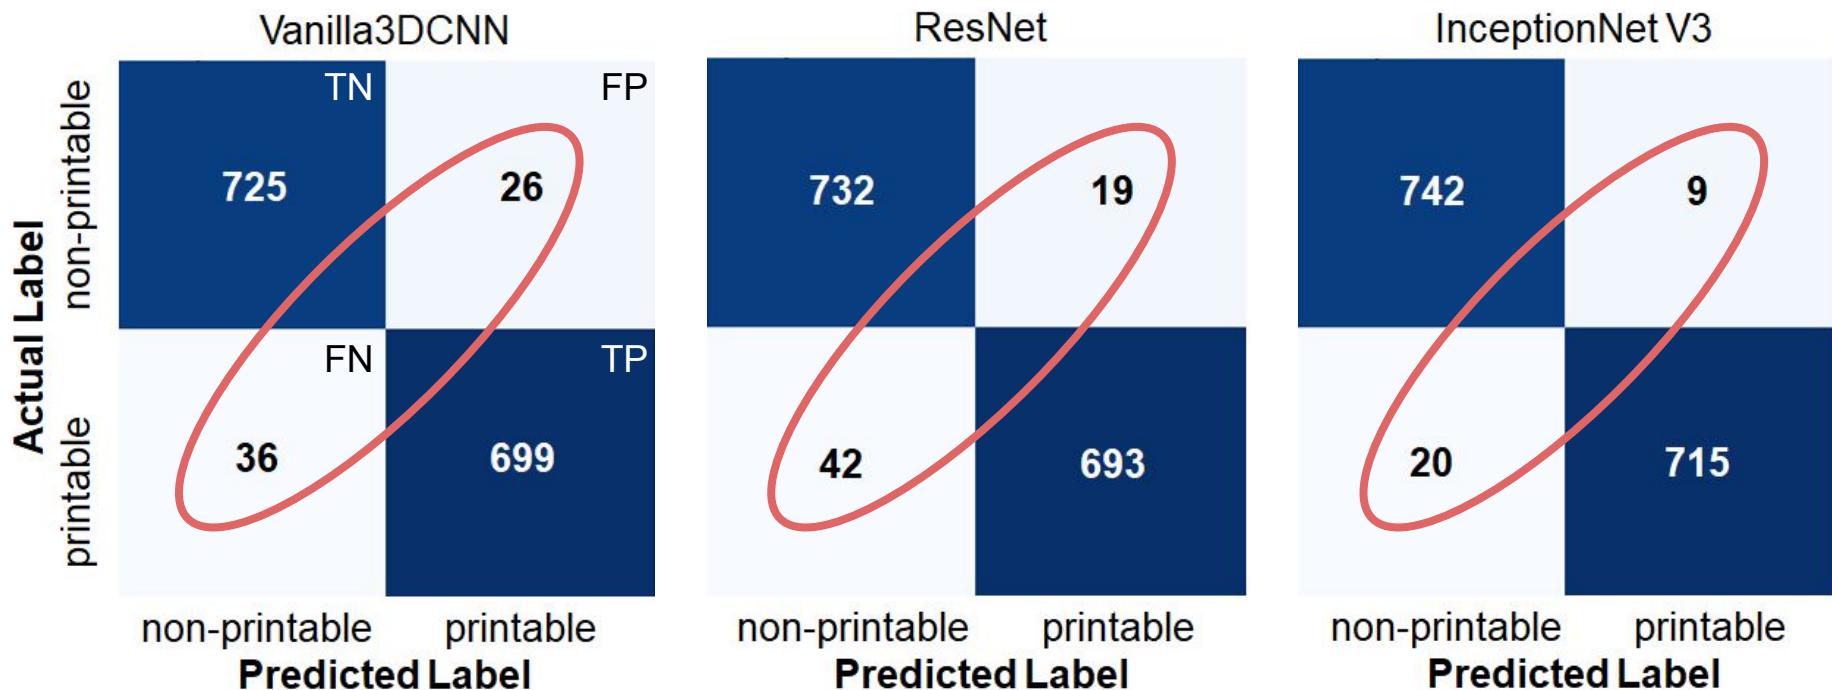
- **Diagonal elements:** correct prediction (TN & TP)
- **Off-diagonal elements:** wrong predictions (FN & FP)

Confusion Matrix



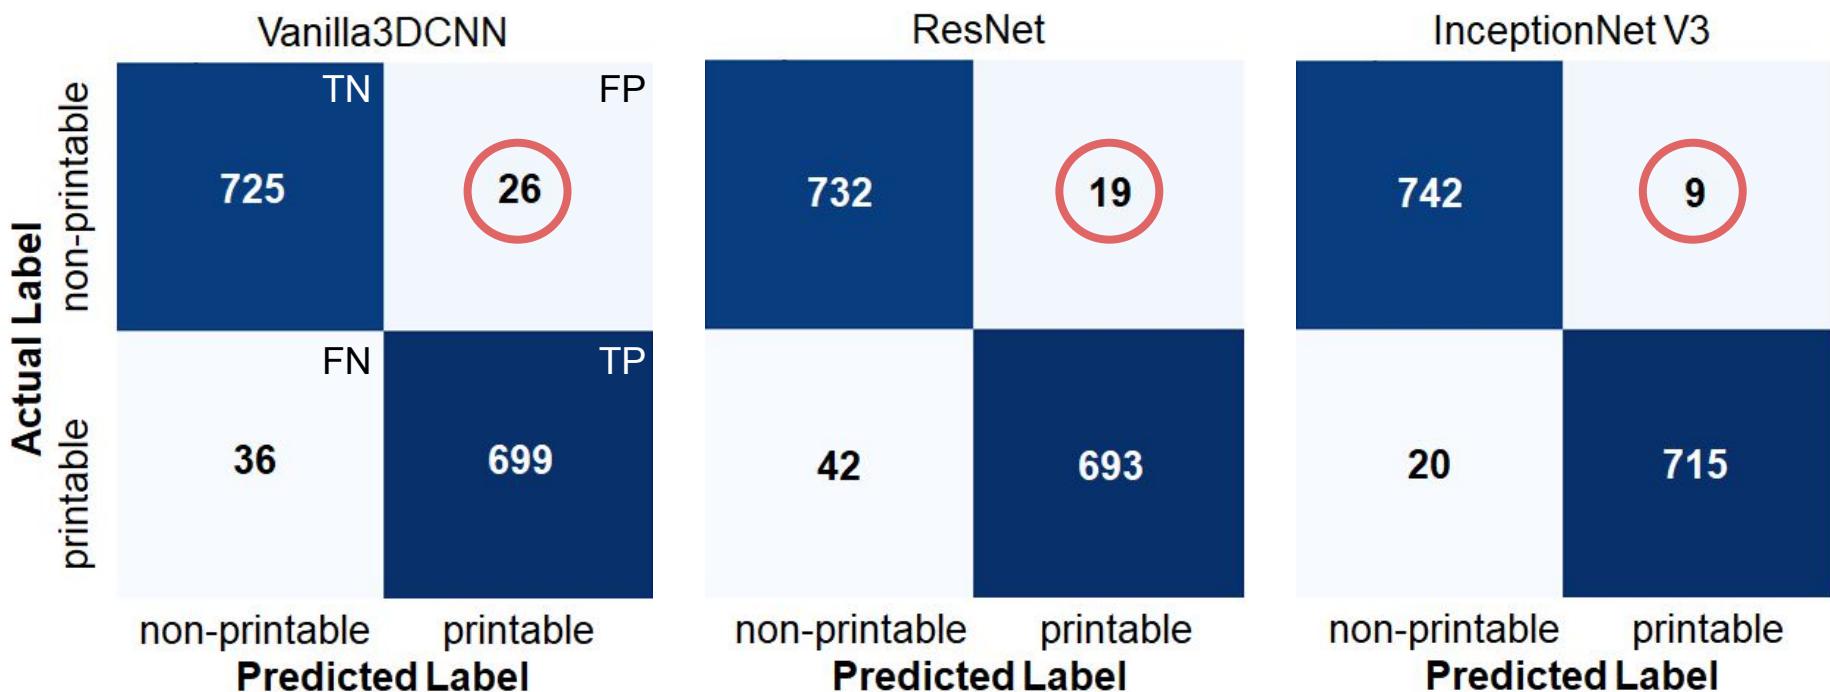
- **True Negatives > True Positives** (for all classifiers)
- **True Negatives:** Method of adding defects works pretty well
- **True Positives:** Assumption of all initial models being printable does not hold

Confusion Matrix



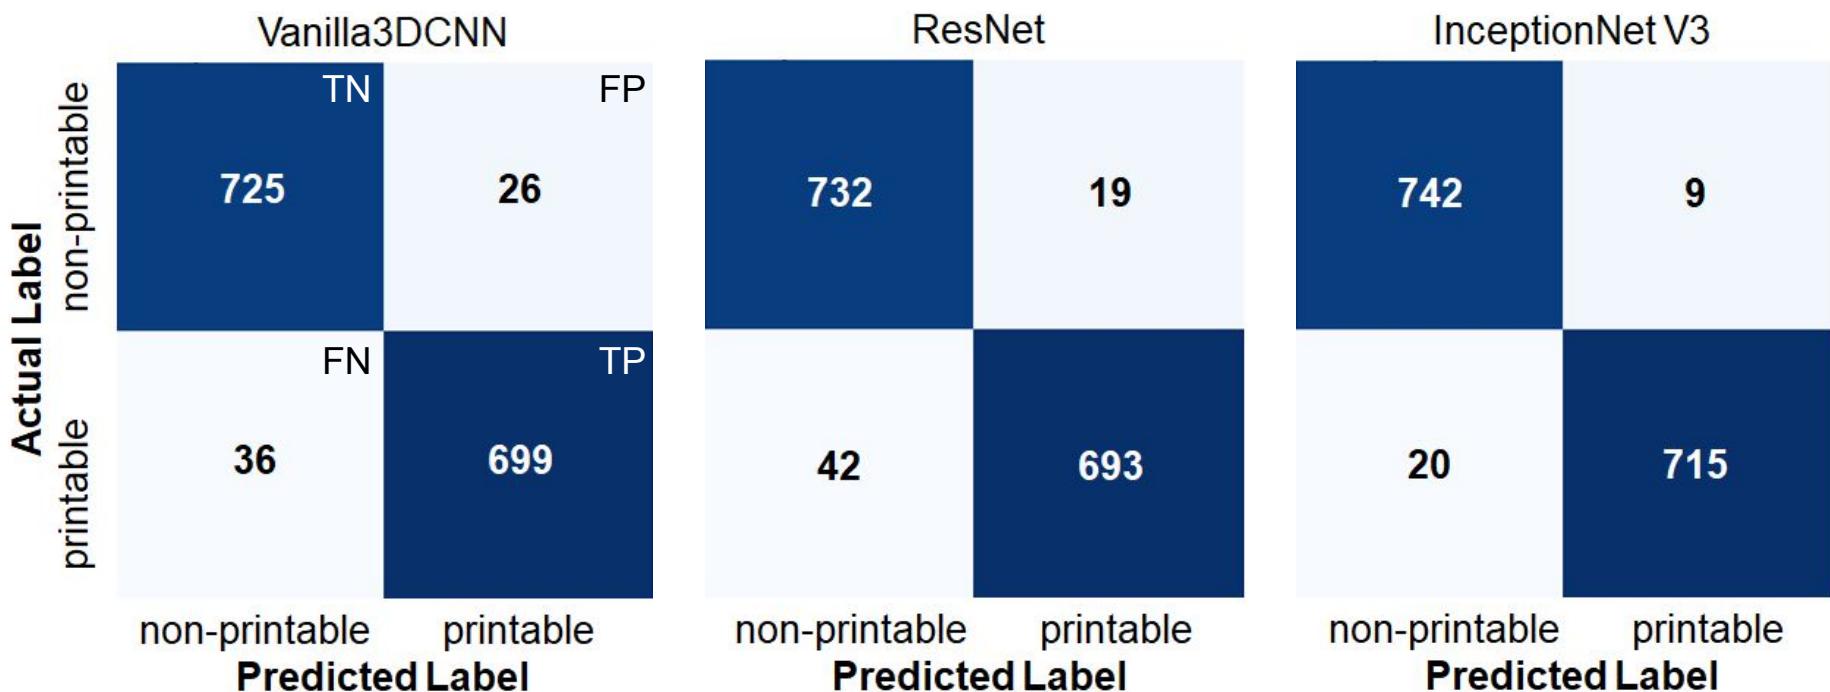
- **Very few off-diagonal elements:** Good overall performance
- **False Negatives > False Positives** (for all classifiers)

Confusion Matrix



- **Avoid:** False Positives (FP) over False Negatives (FN)
- **Additive Manufacturing:** False Positives waste time & material
False Negatives require manual check by engineer

Confirm Results: Accuracy⁽¹⁾ & F1-Score⁽²⁾



Accuracy	F1-Score
95.8 %	0.957

Accuracy	F1-Score
95.9 %	0.959

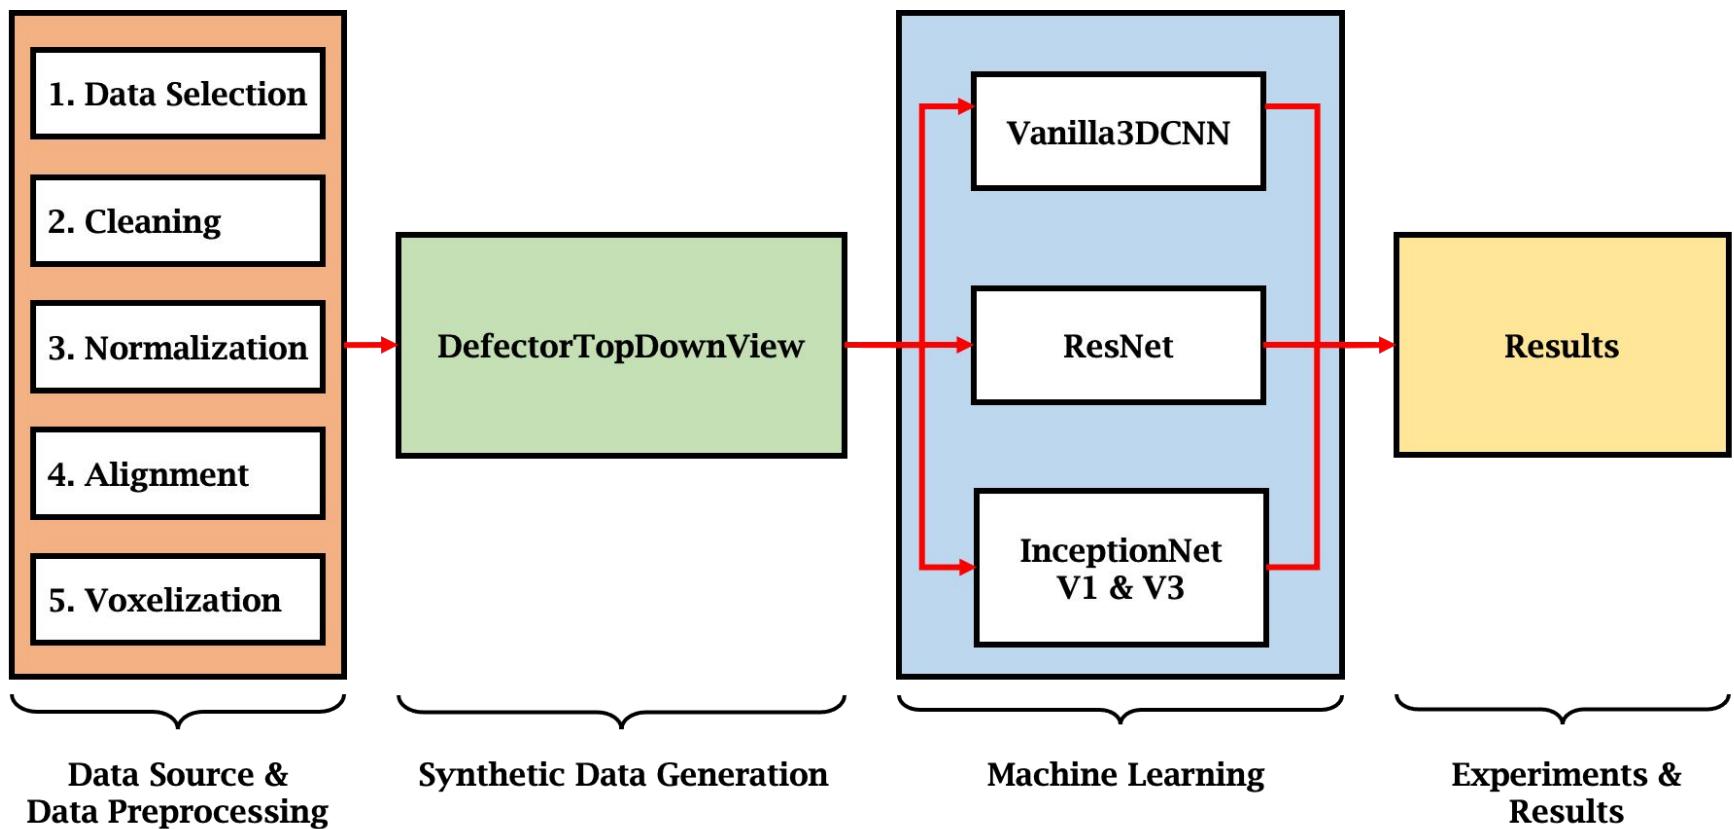
Accuracy	F1-Score
98.1 %	0.980

(1) In this case: balanced accuracy

(2) F1-Score $\in [0, 1]$; harmonic mean of precision⁽³⁾ and recall⁽⁴⁾

(3) Precision: How often is it correct when a positive is predicted: TP/(FP+TP)

(4) Recall: How often is a positive predicted when it actually is positive: TP/(FN+TP) 54



Wrap up

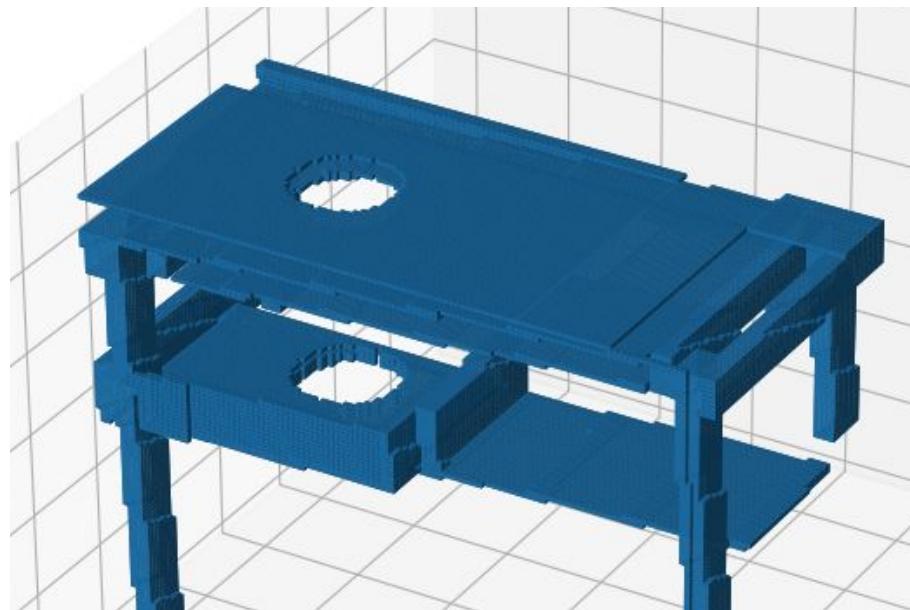
- Good performance of the model using the AMC dataset.
- **Limitations:**
 - Assumption of printability of the models selected
 - Restricted model selection
 - Failure of the defectors in some specific cases

Improvements: Evaluation methods

- Generate an additional test set from another chunk of the **ABC dataset** or from the **Thingi10K dataset**
- **Explainable AI techniques** for better understanding of the performances of the deep learning models.

Improvements: DefectorTopDownView Similarity Check Add-on

- The information given for the tdv could be misleading.
- **Global** and a **local** uniformity **check** for the area that will be removed by the defector.

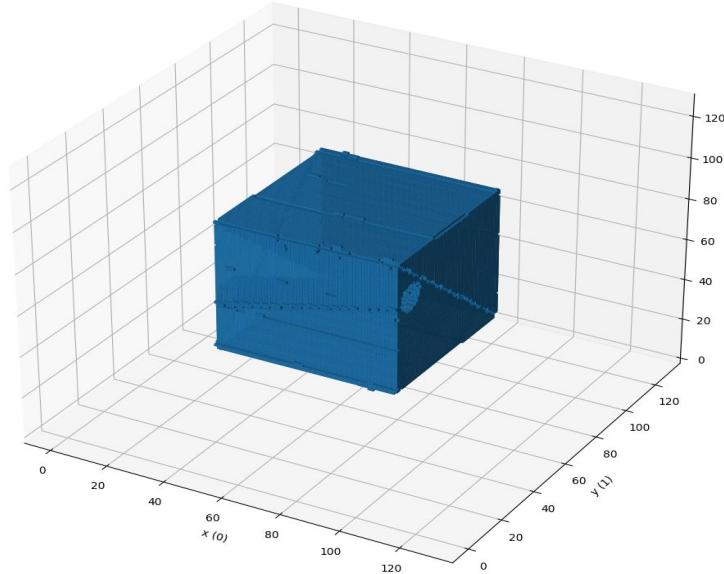


3D model labeled by the defector as manufacturable

Improvements: DefectorTopDownView Rotation Add-on

Insert rotated holes with random angles as a generalization of DefectorTopDownView

- Add a padding to the 3D models
- Rotate the models randomly throughout the x,y and z axes by the angles ϕ_x , ϕ_y , ϕ_z
- Insert the hole using DefectorTopDownView
- Rotate back the model



3D model augmented by a rotated hole

Potential Future Work

The defector is a deterministic algorithm with **limited capabilities**



- Use **self supervision**:
 - Train on original 3D models without artificial defect using an autoencoder
 - Use the encoder as feature extractor and stack final layer for binary classification
 - Train using expert labeled data
- **Few-shot** learning methods.
- Train **generative adversarial networks** (GANs) on a dataset of few complex models that are expert-labeled and generate new models with complex defects.

Thank you for your attention!

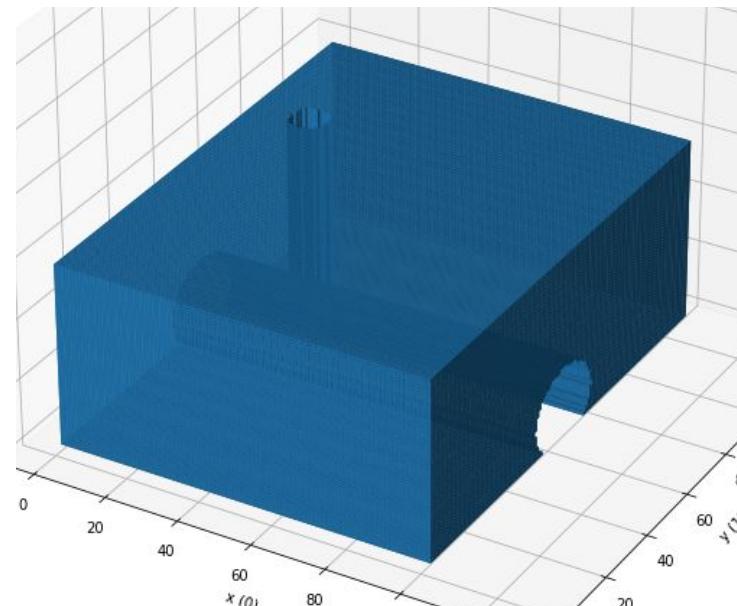
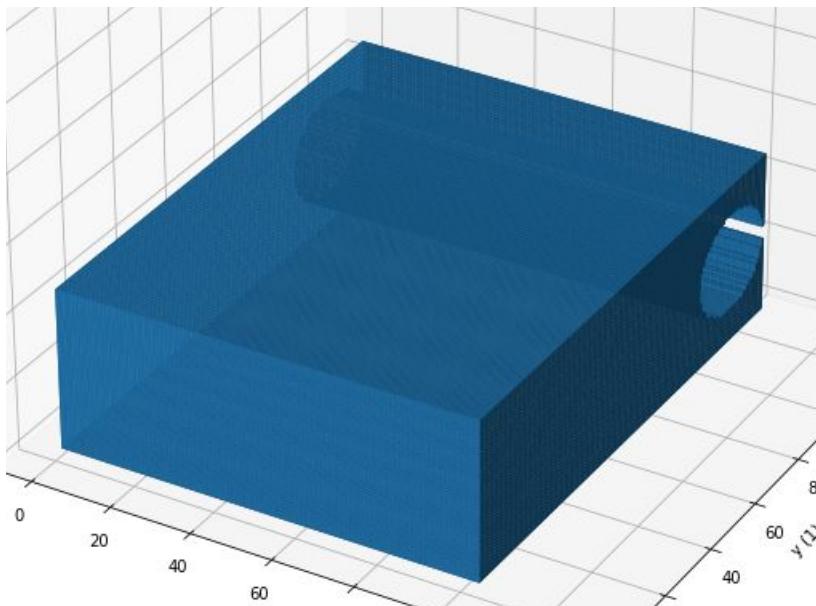
Additional slides

Additional slides

Synthetic Data Generation

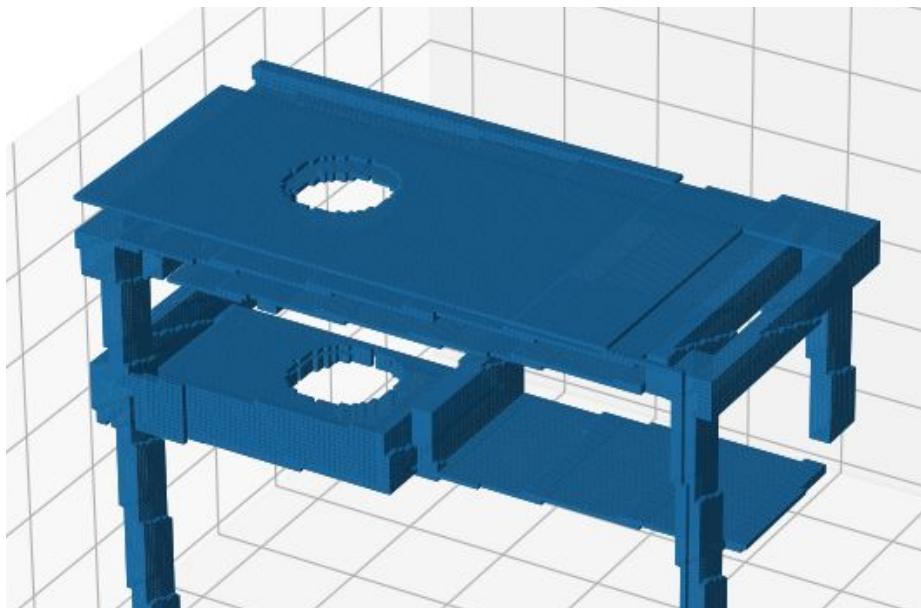
Limitations examples 1/2

Selected 3D models **already contain holes**

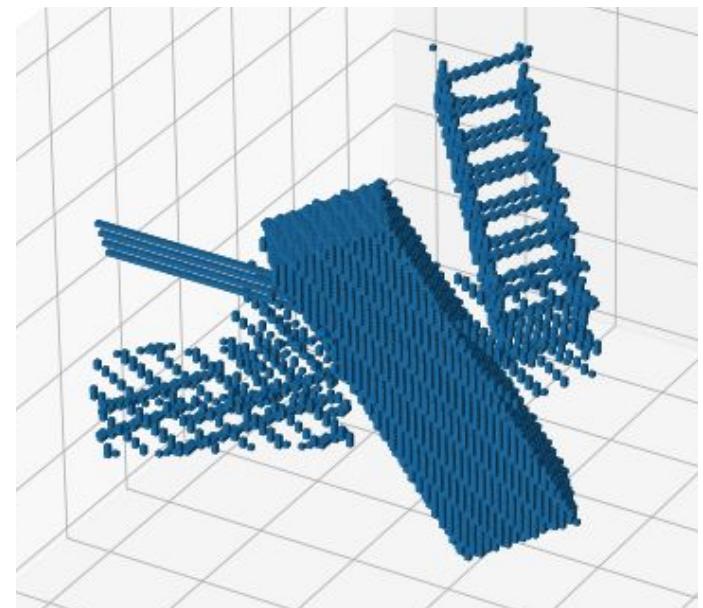


Limitations examples 2/2

Difficult and not detected 3D structure of the model



3D model too complex



Additional slides

Deep Learning

Detailed description of Vanilla3DCNN

Vanilla3DCNN Architecture		
Layer Name	Output Size	CNN - Layers
conv_1	120 x 120 x 120 x 32	9 x 9 x 9/1,32
conv_2	114 x 114 x 114 x 64	7 x 7 x 7/1, 64
Maxpool layer	57 x 57 x 57 x 64	2 x 2 x 2 maxpool,stride 2
conv_3	53 x 53 x 53 x 96	5 x 5 x 5/1, 96
Maxpool layer	26 x 26 x 26 x 96	2 x 2 x 2 maxpool,stride 2
conv_4	24 x 24 x 24 x 128	3 x 3 x 3/1, 128
Maxpool layer	12 x 12 x 12 x 128	2 x 2 x 2 maxpool,stride 2
Average pool layer	11 x 11 x 11 x 128	2 x 2 x 2 maxpool,stride 1
Maxpool layer	1 x 1 x 1 x 128	11 x 11 x 11 maxpool,stride 1
FC_1		32-d fc,ReLU
FC_2		Scalar value, Sigmoid

Detailed description of ResNet

ResNet Architecture		
Layer Name	Output Size	CNN - Layers
conv_1	64 x 64 x 64 x 64	5 x 5 x 5/2,64 (Same)
	32 x 32 x 32 x 64	3 x 3 x 3 maxpool,stride 2
conv_2	32 x 32 x 32 x 64	(3 x 3 x 3/1,64) x 2
		(3 x 3 x 3/1,64) x 2
conv_3	16 x 16 x 16 x 128	(3 x 3 x 3/2,128) x 2
		(3 x 3 x 3/2,128) x 2
conv_4	8 x 8 x 8 x 256	(3 x 3 x 3/2,256) x 2
		(3 x 3 x 3/2,256) x 2
Maxpool layer	4 x 4 x 4 x 256	5 x 5 x 5 maxpool,stride 1
Average pool layer	1 x 1 x 1 x 256	4 x 4 x 4 average pool,stride 1
conv_5	1 x 1 x 1 x 1024	(1 x 1 x 1,256)
FC_1		512-d fc,ReLU
FC_2		128-d fc,ReLU
FC_3		64-d fc,ReLU
FC_4		Scalar value,Sigmoid

Detailed description of InceptionNet V1

InceptionNet_v1 Architecture									
Type	Patch_size/Stride	Output_size	Depth	#1 x 1 x 1	#3 x 3 x 3 reduce	#3 x 3 x 3	#5 x 5 x 5 reduce	#5 x 5 x 5	Pool Proj
Convolution	5 x 5 x 5/1	64 x 64 x 64 x 64	1						
Max pool	2 x 2 x 2/2	32 x 32 x 32 x 64							
Convolution	3 x 3 x 3/1	32 x 32 x 32 x 192	1						
Max pool	2 x 2 x 2/2	16 x 16 x 16 x 192							
Inception(3a)		16 x 16 x 16 x 256	2	64	96	128	16	32	32
max pool	2 x 2 x 2/2	8 x 8 x 8 x 256	0						
Inception (4a)		8 x 8 x 8 x 512	2	192	96	208	16	48	64
Max pool	2 x 2 x 2/2	4 x 4 x 4 x 512	0						
Inception(5a)		4 x 4 x 4 x 1024	2	384	192	384	48	128	128
Maxpool	2 x 2 x 2/2	2 x 2 x 2 x 1024	0						
Inception(6a)		2 x 2 x 2 x 1024	2	384	192	384	48	128	128
Average pool	2 x 2 x 2/1	1 x 1 x 1 x 1024	0						
Dropout	p = 0.4								
FC1		512							
FC2		64							
FC3		1							

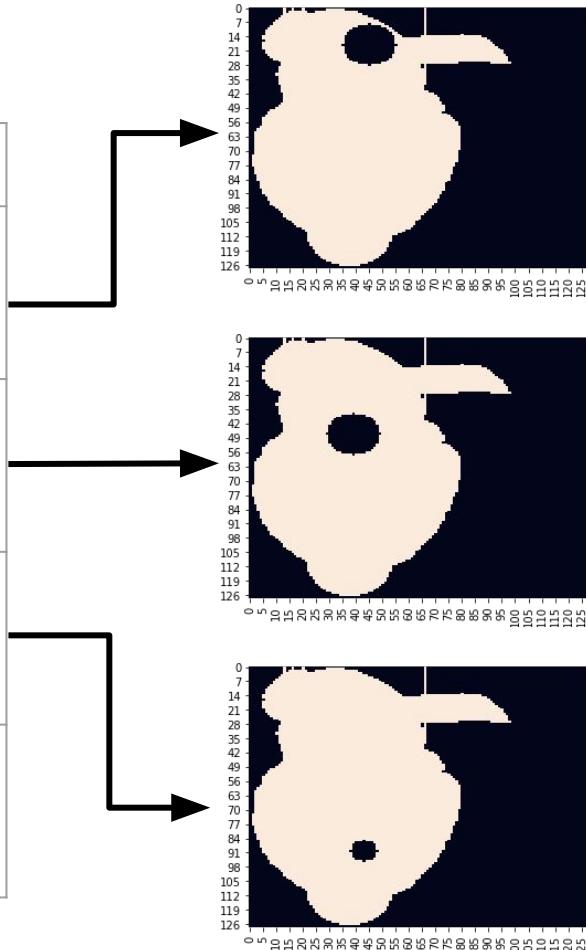
Detailed description of InceptionNet V3

InceptionNet_v3 Architecture										
Type	Patch_size/Stride	Output_size	Depth	#1 x 1 x 1	#3 x 3 x 3 reduce_1	#3 x 3 x 3 reduce_2	#3 x 3 x 3	#3 x 3 x 3 reduce	#3 x 3 x 3	Pool Proj
Convolution	5 x 5 x 5/1	128 x 128 x 128 x 64	1							
Max pool	2 x 2 x 2/2	64 x 64 x 64 x 64								
Convolution	3 x 3 x 3/1	64 x 64 x 64 x 192	1							
Max pool	2 x 2 x 2/2	32 x 32 x 32 x 192								
Inception 3		32 x 32 x 32 x 256	3	32	64	64	128	64	64	32
max pool	2 x 2 x 2/2	16 x 16 x 16 x 256	0							
Inception 4		16 x 16 x 16 x 512	2	128	128	128	128	64	128	128
Max pool	2 x 2 x 2/2	8 x 8 x 8 x 512	0							
Inception 5		4 x 4 x 4 x 1024	2	128	256	256	384	256	384	128
Maxpool	2 x 2 x 2/2	2 x 2 x 2 x 1024	0							
Inception 6		2 x 2 x 2 x 1024	2	128	256	256	384	256	384	128
Average pool	2 x 2 x 2/1	1 x 1 x 1 x 1024	0							
Dropout	p = 0.4									
FC1		512								
FC2		64								
FC3		1								

Additional slides Results & Performance Analysis

False Prediction Analysis⁽¹⁾

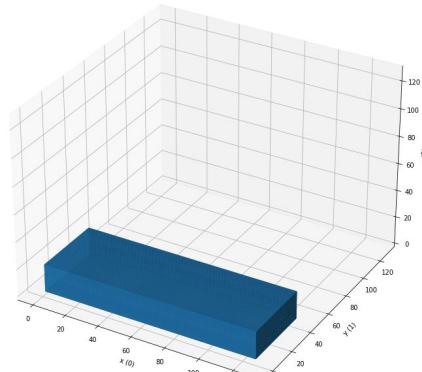
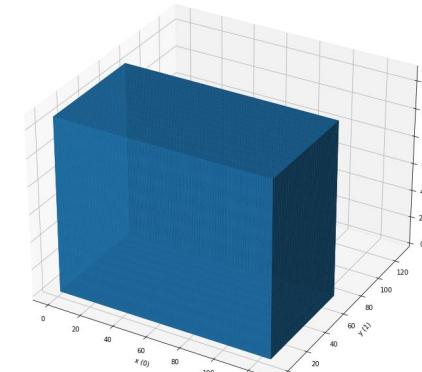
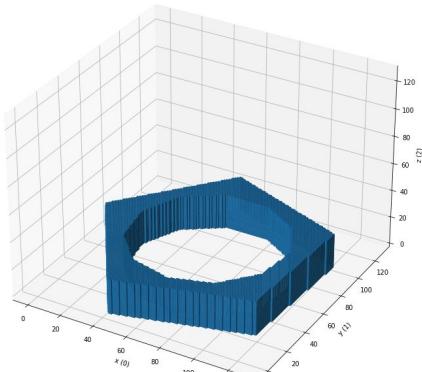
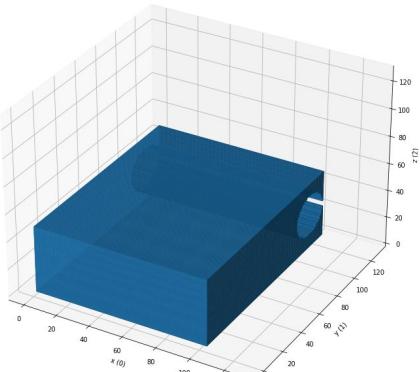
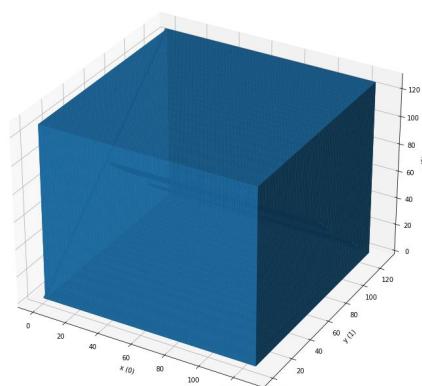
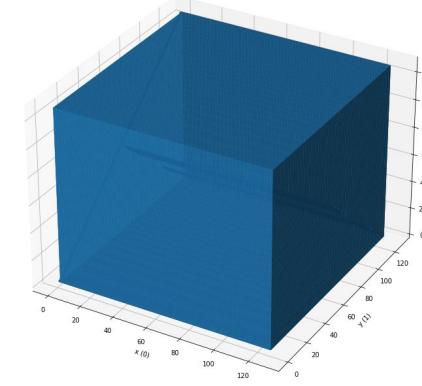
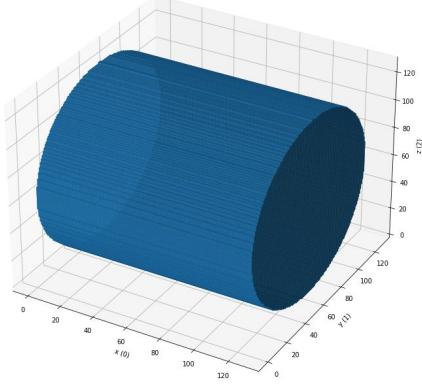
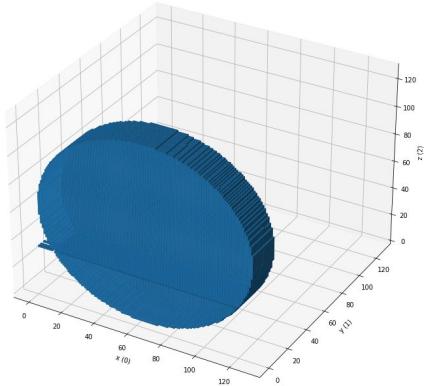
Label	Defect Location	Count
non-printable	border	4
printable	middle	6
non-printable	middle	9
printable	no-defect	10





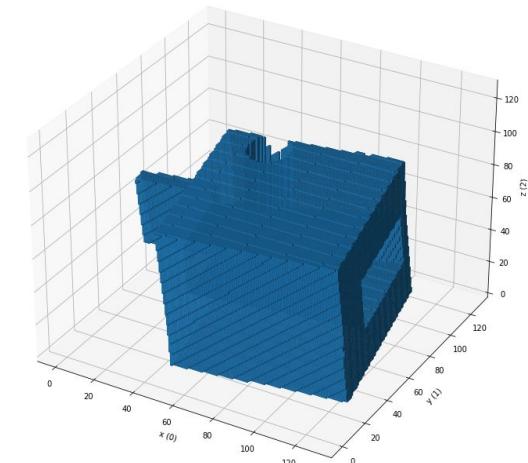
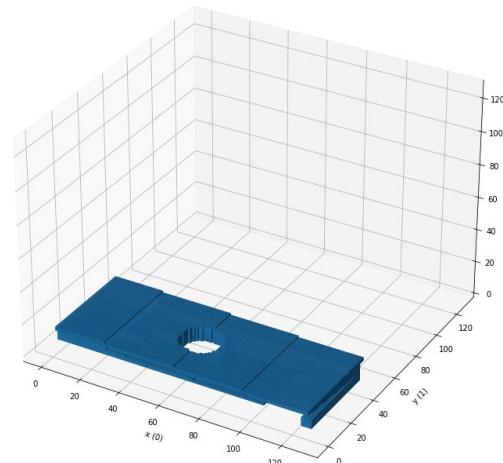
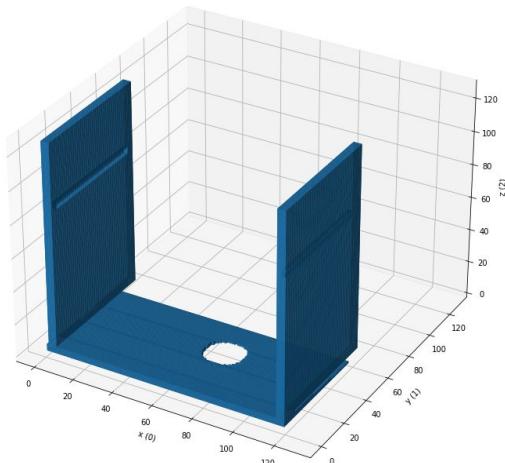
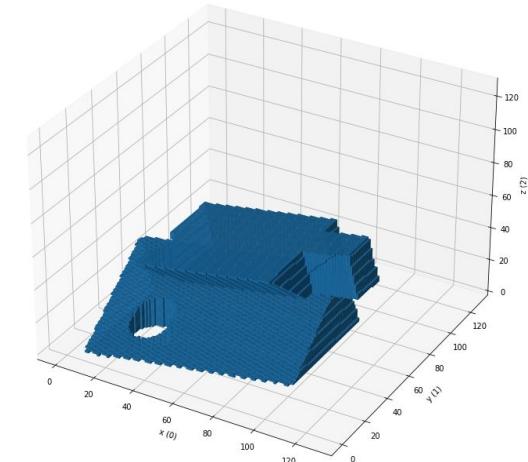
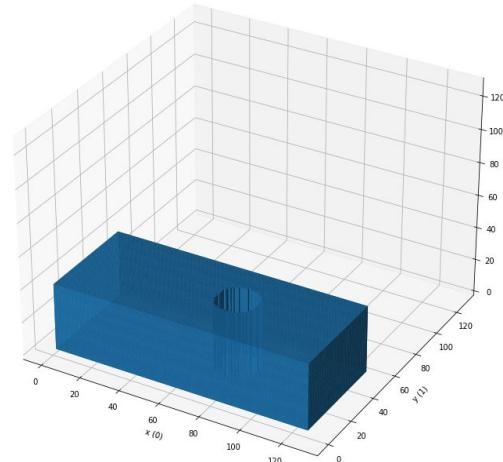
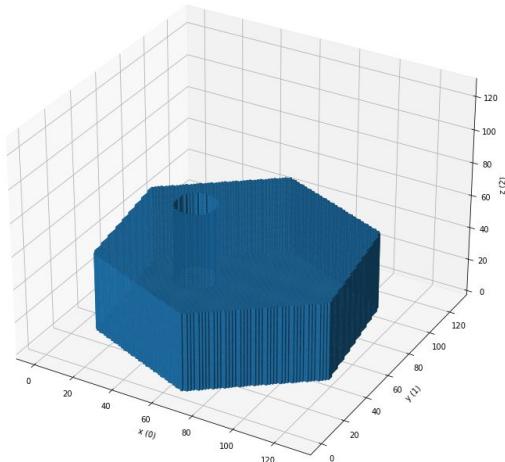
False Prediction Analysis⁽¹⁾

printable - no defect added



False Prediction Analysis⁽¹⁾

printable - middle

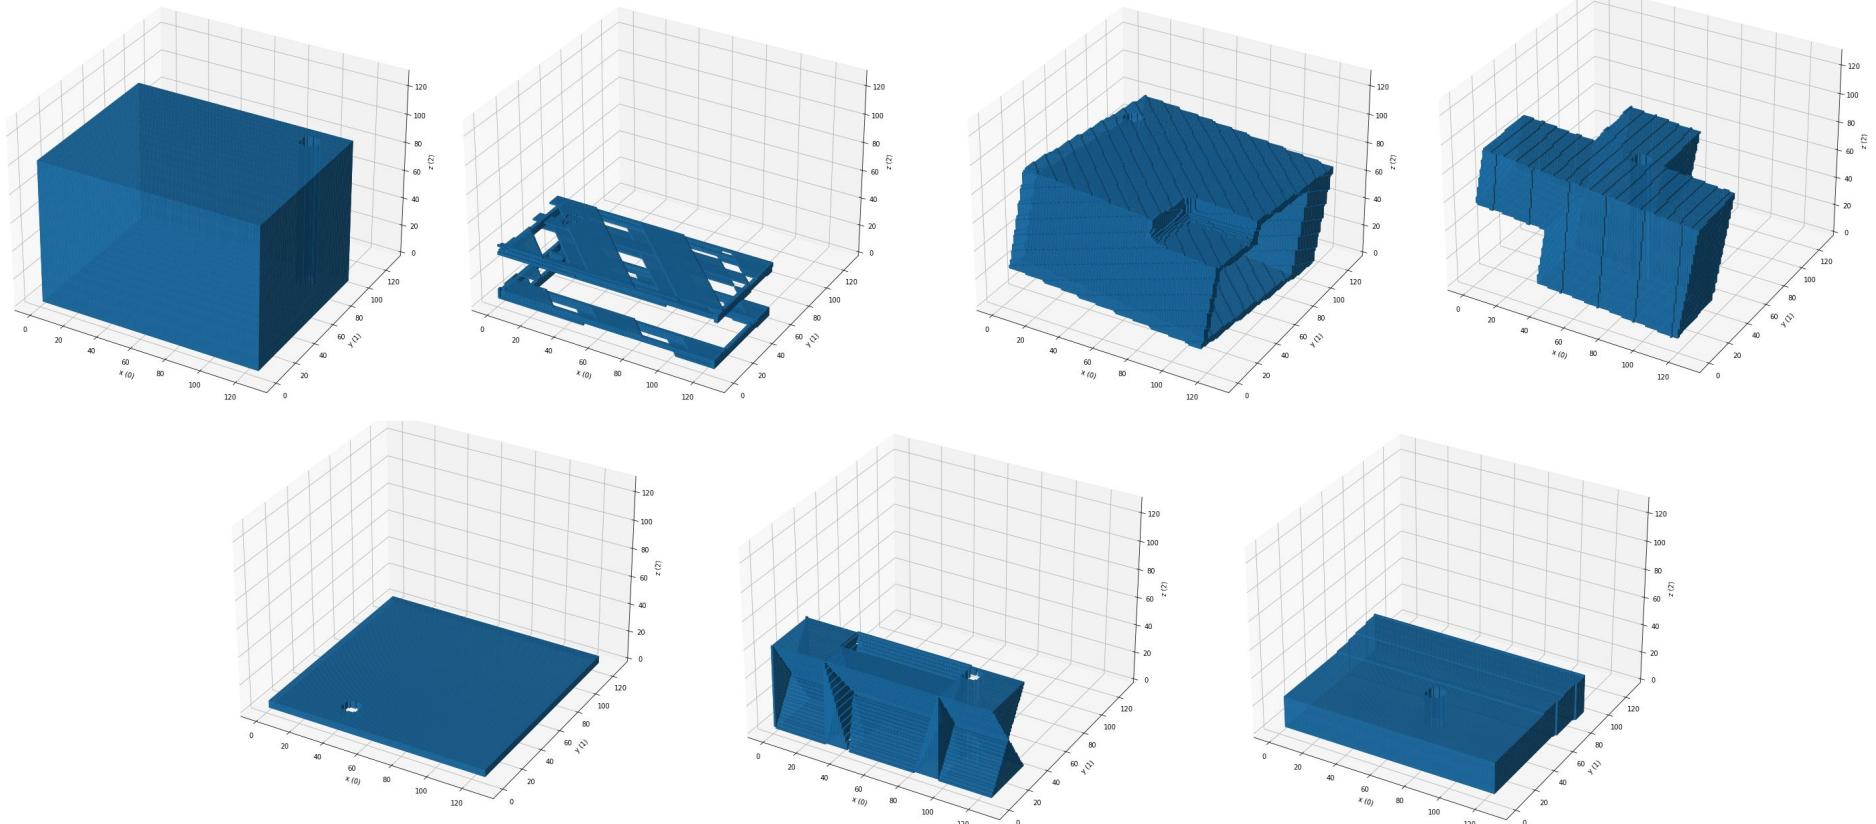
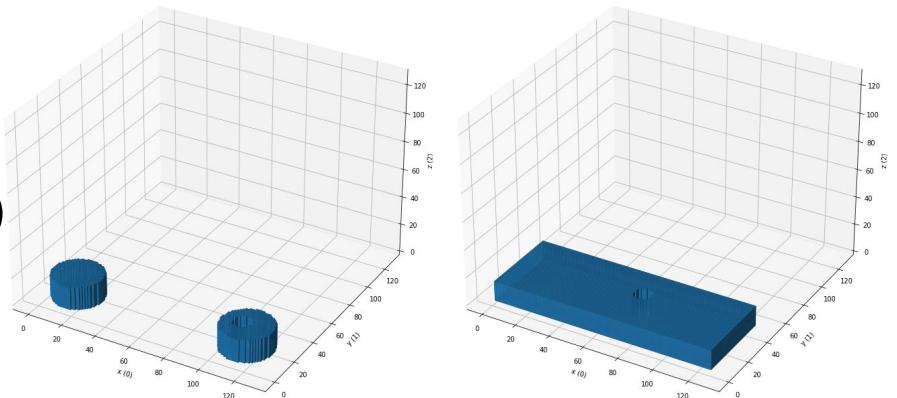


(1) Utilized neural network model: InceptionNet V3



False Prediction Analysis⁽¹⁾

non-printable - middle



False Prediction Analysis⁽¹⁾

non-printable - bord

